

Lecture 7

Linear regression

Julian Reif
Fall 2025

RStudio setup for this lecture

- Log into RStudio on your Amazon EC2 instance
 - Use AMI `FIN550-RStudio` with IAM role `BigDataEC2Role`
- Copy file `lecture-07-tryit.Rmd` from the course Amazon S3 bucket to this folder

```
# Enter this command via RStudio Terminal  
aws s3 cp --recursive s3://bigdata-fin550-reif/lecture-07 ~/fin550/lecture-07
```

Linear regression theory

Recall: machine learning

"All models are wrong, but some are useful"

--- George Box, statistician

- Goal is to model f , the systematic relationship between Y and X

Regression vs classification

- Regression is commonly used when outcome variable is *continuous*
 - Income
 - Spending
 - Height

Goal: model $f(X) = E[Y|X]$, the conditional expectation of Y given X

- By contrast, classification is used when outcome variable is *categorical*
 - Do you attend the University of Illinois? (yes/no)
 - What program are you enrolled in? (MSF/MSFE/MSBA/other)

Goal: model $f(X) = Pr[Y = \text{label}|X]$, the probability that Y is *label*, given X

Linear regression

- Approximates $f(X)$, the conditional expectation function, using a linear function:

$$f(X) = E[Y|X] = \beta_0 + \beta_1 X$$

- Plugging this into $Y = f(X) + \epsilon$ gives the equivalent form of this relationship as:

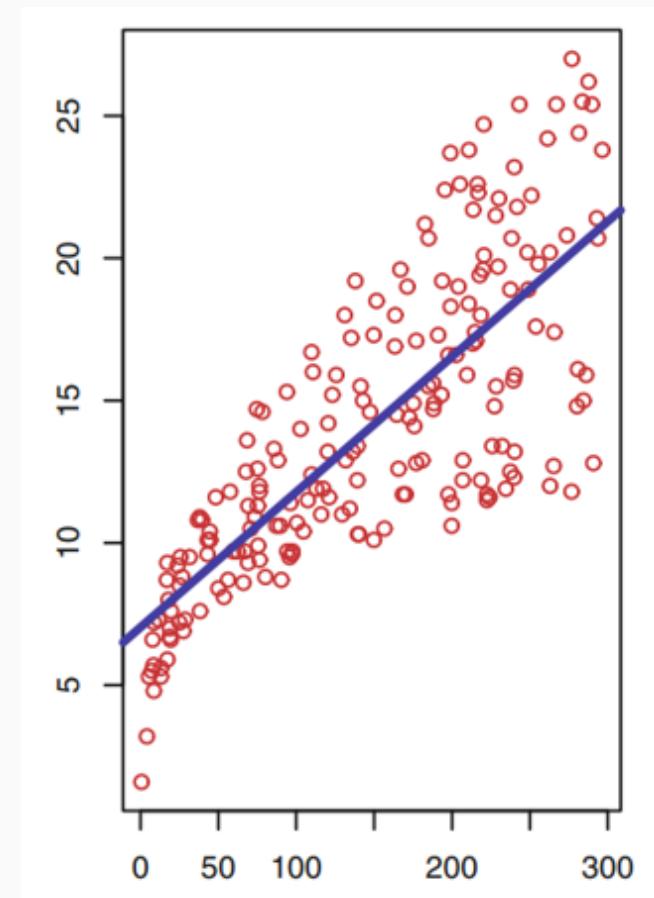
$$Y = \beta_0 + \beta_1 X + \epsilon$$

Univariate linear regression

- Today, we will study univariate linear regression ($X = X_1$)
- Data (x, y) , where x and y are vectors with length equal to n , the number of observations (rows)
- Let x_i and y_i denote the values of x and y for observation i , where $i = 1, \dots, n$

Univariate linear regression

- Can be expressed graphically as a "trendline" drawn through a scatter plot of y versus x



Linear regression

$$Y = \beta_0 + \beta_1 X + \epsilon$$

- β_0 and β_1 are the *parameters* (or *coefficients*) of the model
- In this univariate linear model, β_0 can be called the "intercept" and β_1 the "slope"
- Error term ϵ represents features that affect Y but are unavailable in our data

Our regression will estimate the true (fixed) parameters

- Let $\hat{\beta}_0$ and $\hat{\beta}_1$ be our best estimates of the parameters β_0 and β_1
- Given these parameter estimates, the model's prediction for observation i is

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

- The prediction error is equal to $\hat{\epsilon}_i = y_i - \hat{y}_i$

Linear regression

- Answers prediction questions
 - "What is our best prediction of Y if we only know X ?"
- Can also be used to describe correlations
 - "What is the relationship between X and Y in the data?"
- Later in the course, we will use linear regression for causal inference
 - "If we change X how will Y change?"

Linear regression minimizes RSS

- Linear regression minimizes the square of the prediction error
- Coefficients $\hat{\beta}_0$ and $\hat{\beta}_1$ are chosen to minimize the Residual Sum of Squares (RSS):

$$\begin{aligned} RSS &= \sum_{i=1}^n \hat{\epsilon}_i^2 \\ &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \sum_{i=1}^n \left(y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i) \right)^2 \end{aligned}$$

- Why minimize the **square** and not just the sum?

Linear regression solution

- Minimizing RSS is a standard calculus optimization problem
- Let \bar{y} and \bar{x} be the mean of y and x , respectively
- The solution to this minimization problem is:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

Linear regression in R

Load libraries, turn off scientific notation

```
library(tidyverse)  
library(ggplot2)
```

```
# scipen: a penalty for deciding whether to print fixed or exponential notation.  
# Positive values encourage fixed notation, negative encourage scientific notation  
options(scipen=999)
```

Load and inspect our data

```
housing <- read_csv("lecture-07-housing.csv")
nrow(housing)
names(housing)

# [1] 300
# [1] "PID"                 "Lot_Area"            "House_Style"        "Overall_Qual"
# [5] "Overall_Cond"        "Year_Built"          "Heating_QC"         "Central_Air"
# [9] "Gr_Liv_Area"         "Bedroom_AbvGr"       "Fireplaces"         "Garage_Area"
# [13] "Mo_Sold"              "Yr_Sold"              "SalePrice"           "Basement_Area"
# [17] "Full_Bathroom"        "Half_Bathroom"       "Total_Bathroom"     "Deck_Porch_Area"
# [21] "Age_Sold"              "Season_Sold"          "Garage_Type_2"       "Foundation_2"
# [25] "Masonry_Veneer"        "Lot_Shape_2"          "House_Style2"        "Overall_Qual2"
# [29] "Overall_Cond2"         "Log_Price"            "Bonus"                "score"
```

Ames Iowa Housing Data, 2006-2010

Variable name	Definition
PID	Parcel identification number
Lot_Area	Lot size in square feet
House_Style	Style of dwelling
Overall_Qual	Rates the overall material and finish of the house
Overall_Cond	Rates the overall condition of the house
Year_Built	Original construction date
Heating_QC	Heating quality and condition
Gr_Liv_Area	Above grade (ground) living area square feet
SalePrice	Selling price
Mo_Sold	Month sold
Season_Sold	Season sold

Summarize a few variables

```
summary(housing$SalePrice)
```

```
summary(housing$Mo_Sold)
```

```
#      Min. 1st Qu. Median      Mean 3rd Qu.      Max.  
# 35000 114000 135000 137525 159238 290000  
#      Min. 1st Qu. Median      Mean 3rd Qu.      Max.  
# 1.000 4.000 6.000 5.907 7.000 12.000
```

Build a model of SalePrice, with just an intercept

$$SalePrice = \beta_0 + \epsilon$$

```
lm1 <- lm(SalePrice ~ 1, data = housing) # Outcome variable is SalePrice  
summary(lm1)
```

```
#  
# Call:  
# lm(formula = SalePrice ~ 1, data = housing)  
#  
# Residuals:  
#       Min        1Q    Median        3Q       Max  
# -102525 -23525   -2525   21713  152475  
#  
# Coefficients:  
#                 Estimate Std. Error t value     Pr(>|t|)  
# (Intercept)  137525      2172   63.31 <0.000000000000002 ***  
# ---  
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
#
```

Model with only an intercept

```
# lm1 is a linear model "object"  
  
# Extract the value of the alpha coefficient from this object  
lm1$coefficients  
  
# (Intercept)  
#       137524.9  
  
# Note: in this intercept model, the coefficient is just the mean  
mean(housing$SalePrice)  
  
# [1] 137524.9
```

Model with a numeric (continuous) predictor

- Estimate a model where sale price is a function of above-ground living area

$$SalePrice = \beta_0 + \beta_1 \text{Gr_Liv_Area} + \epsilon$$

- Above-ground living area is measured in square feet

```
summary(housing$Gr_Liv_Area)
```

```
#      Min. 1st Qu. Median      Mean 3rd Qu.      Max.  
#      334     952    1135     1131    1337     1500
```

Try it: Model with a numeric (continuous) predictor

```
lm2 <- lm(SalePrice ~ Gr_Liv_Area, data = housing)  
summary(lm2)
```

Try it: Model with a numeric (continuous) predictor

```
lm2 <- lm(SalePrice ~ Gr_Liv_Area, data = housing)
summary(lm2)

#
# Call:
# lm(formula = SalePrice ~ Gr_Liv_Area, data = housing)
#
# Residuals:
#   Min     1Q Median     3Q    Max
# -77582 -20114     358 19208 122153
#
# Coefficients:
#             Estimate Std. Error t value    Pr(>|t|)    
# (Intercept) 18583.434   8213.438   2.263 0.0244 *  
# Gr_Liv_Area 105.189      7.115 14.784 <0.000000000000002 *** 
# ---
# Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#
# Residual standard error: 28620 on 298 degrees of freedom
# Multiple R-squared:  0.4231,    Adjusted R-squared:  0.4212 
# F-statistic: 218.6 on 1 and 298 DF,  p-value: < 0.000000000000022
```

- Parameter estimate for the intercept is 18583.434
- Parameter estimate for the slope of Gr_Liv_Area is 105.189
- Regression equation is $SalePrice=18583.434 + 105.189*Gr_Liv_Area$
- Model indicates that each additional square foot of Gr_Liv_Area is associated with about \$105 higher sale price

Model with a numeric (continuous) predictor

```
# The lm2 object includes predicted values ("y hat")
```

```
head(lm2$fitted.values)
```

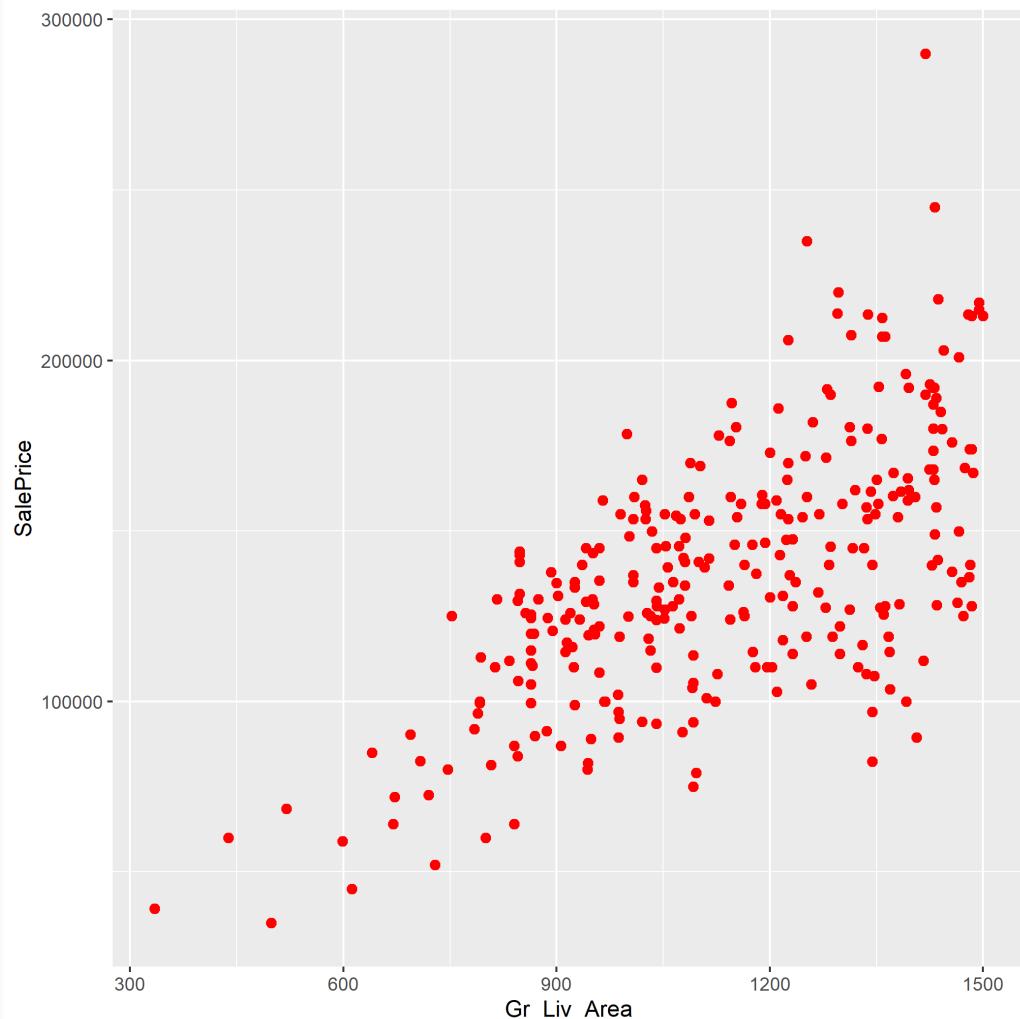
```
#          1         2         3         4         5         6  
# 159326.34 153225.38 109466.75 139024.86 169003.73  97685.58
```

Try it: scatter plot plus trendline

```
# Plot SalePrice (y-axis) vs Gr_Liv_Area (x-axis)
# Make the scatterpoints red, with size=2
ggplot(housing, aes( )) +
  geom_point()
```

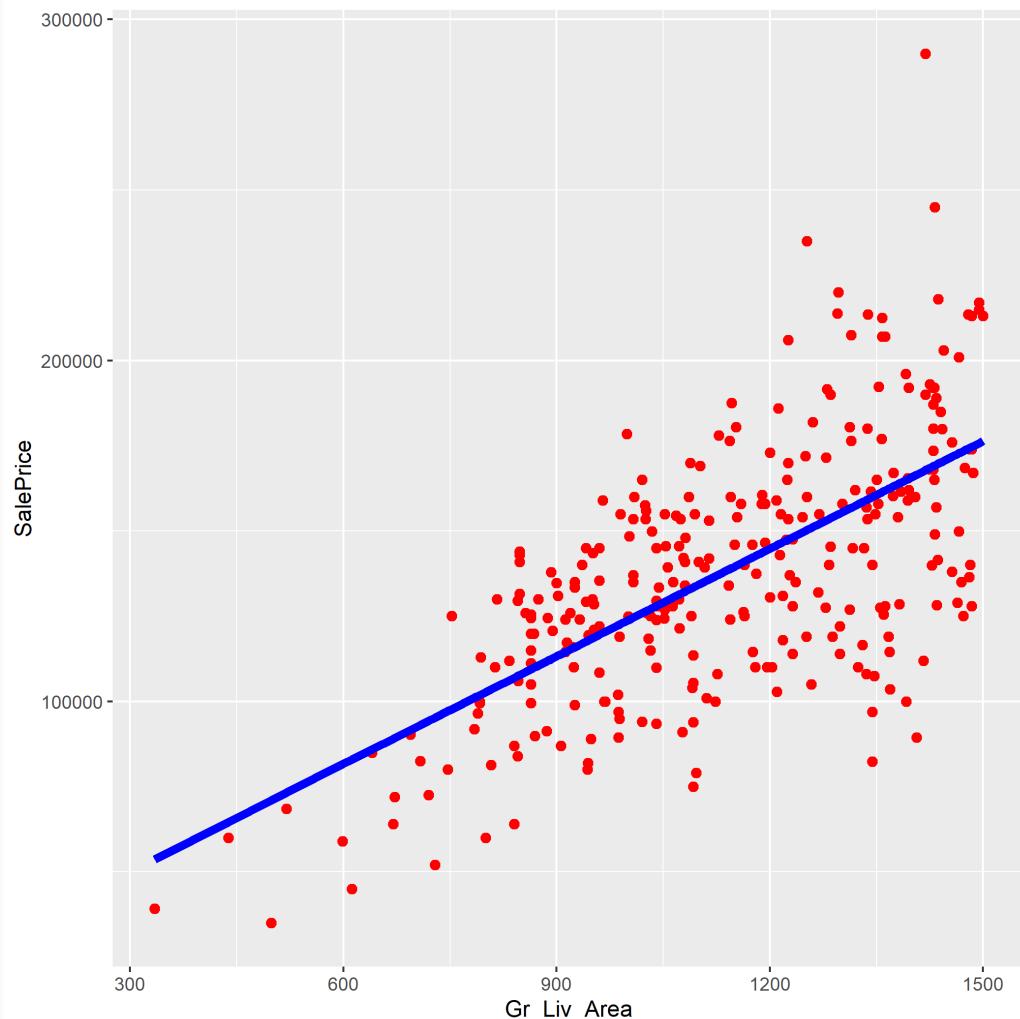
Try it: scatter plot plus trendline

```
# Plot SalePrice (y-axis) vs Gr_Liv_Area (x-axis)
# Make the scatterpoints red, with size=2
ggplot(housing, aes(x = Gr_Liv_Area, y = SalePrice)) +
  geom_point(color="red", size=2)
```



Try it: scatter plot plus trendline

```
# Plot SalePrice (y-axis) vs Gr_Liv_Area (x-axis)
# Make the scatterpoints red, with size=2
ggplot(housing, aes(x = Gr_Liv_Area, y = SalePrice)) +
  geom_point(color="red", size=2) +
  geom_line(aes(x = Gr_Liv_Area, y = lm2$fitted.values),
            color="blue", linewidth=2)
```



Model with categorical predictor

- Does the season when the house is sold matter for the sale price?
- Season_Sold is a categorical predictor

```
table(housing$Season_Sold)
```

```
#  
#   1   2   3   4  
# 35 100 130 35
```

- Does treating it as a continuous variable make sense?

$$SalePrice = \beta_0 + \beta_1 \text{Season_Sold} + \epsilon$$

- No!

Model with categorical predictor

- Categorical predictors must be split up into their respective categories
- Create a set of 3 (= 4 - 1) "dummy" variables that are equal to either 0 or 1
 - dummy2: equal to 1 if Season_Sold==2, 0 otherwise
 - dummy3: equal to 1 if Season_Sold==3, 0 otherwise
 - dummy4: equal to 1 if Season_Sold==4, 0 otherwise
- What happened to Season_Sold==1 ?
 - It is an "omitted" (baseline) category

Model with categorical predictor

- The function `as.factor()` automatically creates dummy variables for us

```
lm(SalePrice ~ as.factor(Season_Sold), data = housing)

#
# Call:
# lm(formula = SalePrice ~ as.factor(Season_Sold), data = housing)
#
# Coefficients:
#             (Intercept)  as.factor(Season_Sold)2  as.factor(Season_Sold)3
#                   120917                  19962                  16493
# as.factor(Season_Sold)4
#                   24060
```

`SalePrice = 120917 + 19962*dummy2 + 16493*dummy3 + 24060*dummy4`

Model with categorical predictor

`SalePrice = 120917 + 19962*dummy2 + 16493*dummy3 + 24060*dummy4`

- Predicted price for house sold in "omitted" season 1 (winter): 120917
- Predicted price for house sold in season 2 (spring): $120917 + 19962 = 140879$
- Predicted price for house sold in season 3 (summer): $120917 + 16493 = 137410$
- Predicted price for house sold in season 4 (fall): $120917 + 24060 = 144977$
- Note: these predictions are just the `SalePrice` means for these groups

Categorical regressions are means

```
SalePrice = 120917 + 19962*dummy2 + 16493*dummy3 + 24060*dummy4
```

```
housing %>%
  group_by(Season_Sold) %>%
  summarize(SalePrice_mean = mean(SalePrice))
```

```
# # A tibble: 4 × 2
#   Season_Sold SalePrice_mean
#       <dbl>          <dbl>
# 1 1           120917.
# 2 2           140879.
# 3 3           137410.
# 4 4           144977.
```

Categorical regressions are means

- Regressions with a categorical predictor always yield means
- Note: not necessarily true if you also include other predictors

```
# Does not report simple means
lm(SalePrice ~ Gr_Liv_Area + as.factor(Season_Sold), data = housing)

#
# Call:
# lm(formula = SalePrice ~ Gr_Liv_Area + as.factor(Season_Sold),
#     data = housing)
#
# Coefficients:
#             (Intercept)          Gr_Liv_Area  as.factor(Season_Sold)2
#                   13395.4              103.7            6813.0
# as.factor(Season_Sold)3  as.factor(Season_Sold)4
#                   8185.0              8593.2
```

What happens if the regressor has non-numeric values?

```
# "Ex" = "Excellent", "Gd" = "Good", "TA" = "Typical", "Fa" = "Fair"
typeof(housing$Heating_QC)
table(housing$Heating_QC)
lm(SalePrice ~ Heating_QC, data = housing)

# [1] "character"
#
#   Ex   Fa   Gd   TA
# 107  16  58 119
#
# Call:
# lm(formula = SalePrice ~ Heating_QC, data = housing)
#
# Coefficients:
# (Intercept) Heating_QCFa  Heating_QCGd  Heating_QCTA
#           154919        -57800       -24075       -24346
```

Creating dummy variables by hand

```
# First 6 values  
head(housing$Heating_QC)  
  
# Create dummies using model.matrix. 0 specifies no intercept  
head(model.matrix(~ 0 + Heating_QC, data = housing))
```

```
# [1] "Ex" "Ex" "TA" "Ex" "Ex" "TA"  
#   Heating_QCEx Heating_QCFa Heating_QCGd Heating_QCTA  
# 1       1         0         0         0  
# 2       1         0         0         0  
# 3       0         0         0         1  
# 4       1         0         0         0  
# 5       1         0         0         0  
# 6       0         0         0         1
```

Creating dummy variables by hand

```
# Turn the matrix into a data frame  
Heating_QC_dummy <- as.data.frame(model.matrix(~ 0 + Heating_QC, data = housing))  
head(Heating_QC_dummy)
```

	Heating_QCEx	Heating_QCFa	Heating_QCGd	Heating_QCTA
# 1	1	0	0	0
# 2	1	0	0	0
# 3	0	0	0	1
# 4	1	0	0	0
# 5	1	0	0	0
# 6	0	0	0	1

Creating dummy variables by hand

```
# Add to housing dataset, excluding the first dummy variable  
housing <- cbind(housing, Heating_QC_dummy[, -1])  
  
# fit a linear regression of Sale Price on three dummies  
lm(SalePrice ~ Heating_QCFa + Heating_QCGd + Heating_QCTA, data = housing)
```

```
#  
# Call:  
# lm(formula = SalePrice ~ Heating_QCFa + Heating_QCGd + Heating_QCTA,  
#     data = housing)  
#  
# Coefficients:  
# (Intercept) Heating_QCFa Heating_QCGd Heating_QCTA  
# 154919      -57800       -24075      -24346
```

Summary

- Regression is a popular prediction method for continuous outcome variables
- Estimate regression by minimizing the Residual Sum of Squares
- Categorical predictors must be turned into "dummy" variables
- Lab-07 due Sunday at 11:59pm
- Make sure to **stop your instance** when you are done working