

Lecture 10

Cross-validation

Julian Reif
Fall 2025

RStudio setup for this lecture

- Log into RStudio on your Amazon EC2 instance
 - Use AMI **FIN550-RStudio** with IAM role **BigDataEC2Role**

Enter this command using RStudio Terminal

```
aws s3 cp --recursive s3://bigdata-fin550-reif/lecture-10 ~/fin550/lecture-10
```

Statistics review

R-squared is a measure of fit that ranges from 0 to 1

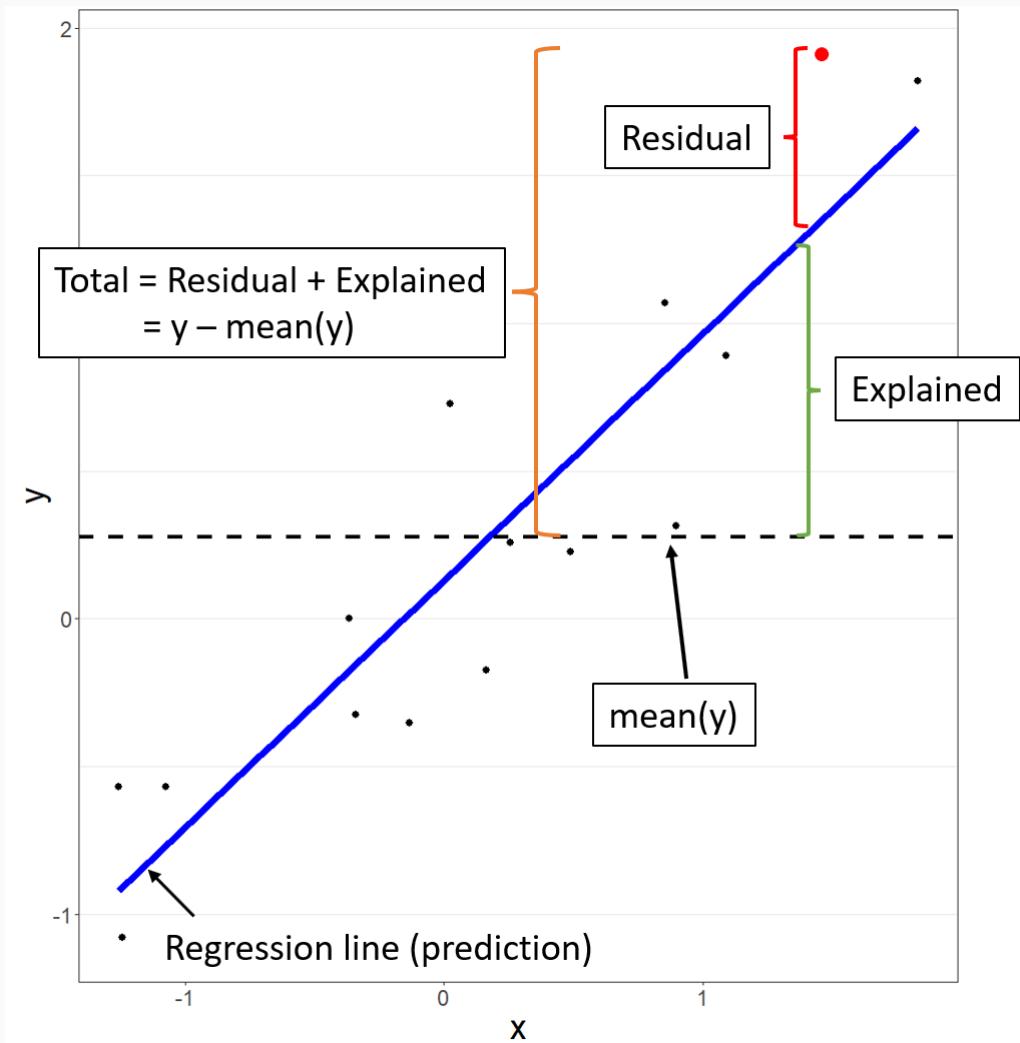
$$\begin{aligned} R^2 &= ESS/TSS \\ &= 1 - RSS/TSS \end{aligned}$$

where:

- ESS: explained sum of squares
- TSS: total sum of squares
- RSS: residual sum of squares (squared prediction error)
- R^2 ranges from 0 to 1
 - $R^2 = 0$: model explains 0% of the variation in Y
 - $R^2 = 1$: model explains 100% of the variation in Y

Illustration of R-squared

- A high R^2 means:
 - Residuals are small
 - Explained variation is high



Linear regression and R-squared

- Recall: linear regression minimizes RSS (squared prediction error)
- Thus, $R^2 = 1 - RSS/TSS$ increases every time a predictor is added to a regression
- When the number of predictors equals the number of observations, $R^2 = 1!$
 - This is an example of "overfitting" your data

Flip a virtual coin 10 times

```
library(tidyverse)
set.seed(26) # Set seed for replicability

flips <- tibble(
  toss = (1:10),
  heads = sample(0:1,
                size = length(toss),
                replace = TRUE)
)

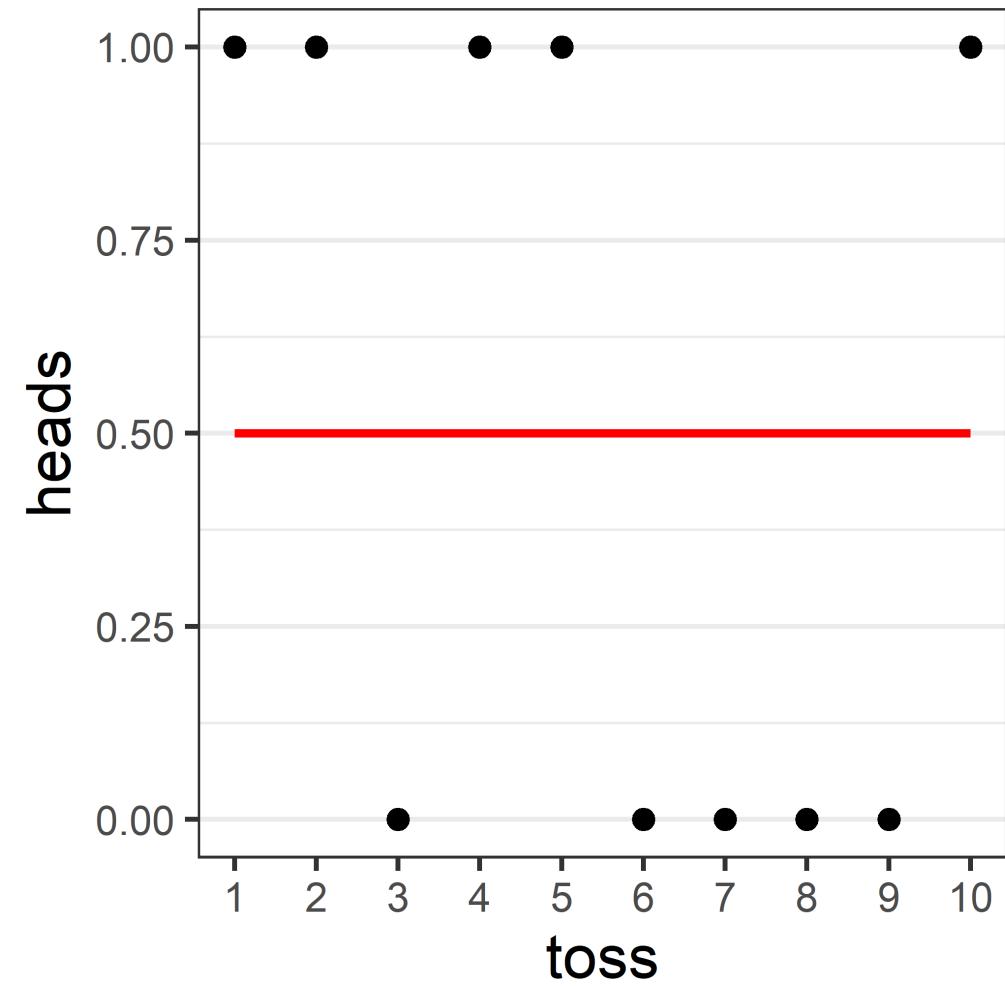
flips
```

#	toss	heads
# 1	1	1
# 2	2	1
# 3	3	0
# 4	4	1
# 5	5	1
# 6	6	0
# 7	7	0
# 8	8	0
# 9	9	0
# 10	10	1

Model 1: no predictors

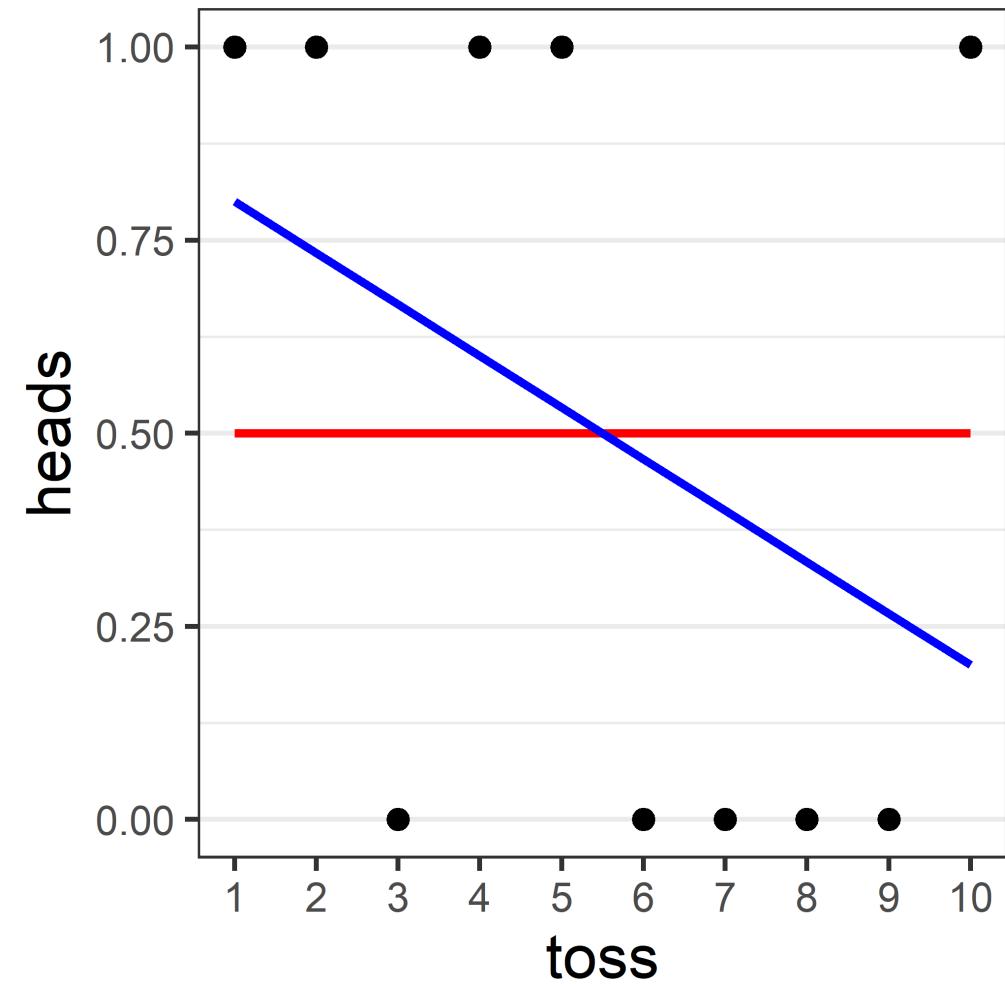
```
# Model 1: no predictors (intercept only)
fit1 <- lm(heads ~ 1,
            data = flips)

# Plot predicted versus actual outcomes
gg <- flips %>%
  ggplot(aes(x = toss, y = heads)) +
  geom_point(size=5) +
  geom_line(aes(y=fit1$fitted.values),
            color="red", linewidth=2) +
  scale_x_continuous(breaks = c(1:10))
gg
```



Model 2: toss order as predictor

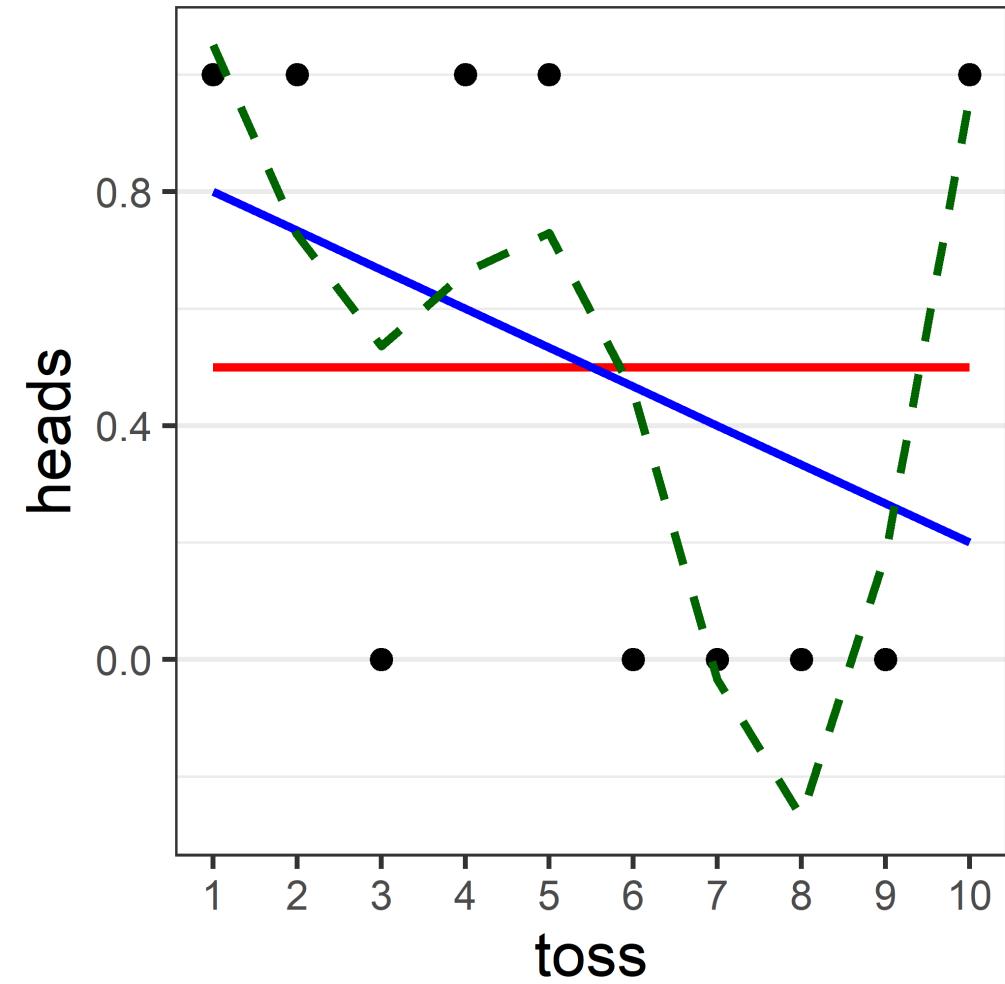
```
# Model 2: toss order as predictor  
fit2 <- lm(heads ~ toss,  
           data = flips)  
  
# Plot predicted versus actual outcomes  
gg <- gg +  
  geom_line(aes(y=fit2$fitted.values),  
            color="blue", linewidth=2)  
gg
```



Model 3: polynomial in toss order (degree 6)

```
# Model 3: toss order poly(6) as predictor
fit3 <- lm(heads ~ poly(toss, 6),
            data = flips)

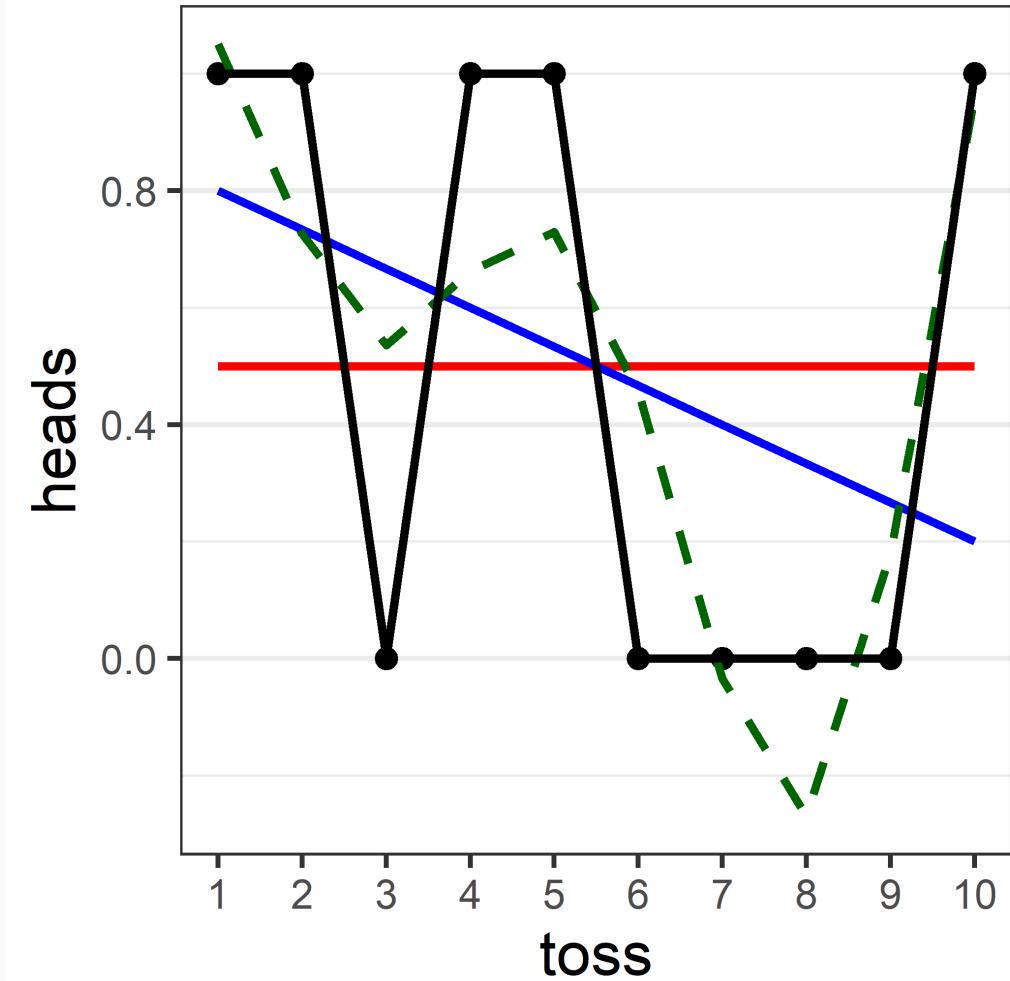
# Plot predicted versus actual outcomes
gg <- gg +
  geom_line(aes(y=fit3$fitted.values),
            color="darkgreen",
            linetype="dashed",
            linewidth=2)
gg
```



Model 4: polynomial in toss order (degree 9)

```
# Model 4: toss order poly(9) as predictor
fit4 <- lm(heads ~ poly(toss, 9),
            data = flips)

# Plot predicted versus actual outcomes
gg <- gg +
  geom_line(aes(y=fit4$fitted.values),
            color="black", linewidth=2)
gg
```



R-squared and MSE for the four models

```
models <- list(fit1, fit2, fit3, fit4)

# Loop over each model and report its MSE & R^2

for(x in 1:4) {

  r2 <- round(summary(models[[x]])$r.squared, 3)

  p <- predict.lm(models[[x]])

  MSE <- round(mean((flips$heads - p)^2), 3)

  print(paste("R2 for model", x, "is:", r2))
  print(paste("MSE for model", x, "is:", MSE))
  cat("\n")
}
```

```
# [1] "R2 for model 1 is: 0"
# [1] "MSE for model 1 is: 0.25"
#
# [1] "R2 for model 2 is: 0.147"
# [1] "MSE for model 2 is: 0.213"
#
# [1] "R2 for model 3 is: 0.652"
# [1] "MSE for model 3 is: 0.087"
#
# [1] "R2 for model 4 is: 1"
# [1] "MSE for model 4 is: 0"
```

Cross-validation

Training data vs test data

- Training data: used to estimate the model (\hat{f})
- Test data: used to evaluate the model
- A model that fits the training data well might not fit the test data well

Create test data for coin-flipping example

```
set.seed(25) # Note: different seed from training data

flips_test_data <- tibble(
  toss = (1:10),
  heads = sample(0:1,
                 size = length(toss),
                 replace = TRUE)
)
```

Compute mean-squared error (MSE) using the test data

```
models <- list(fit1, fit2, fit3, fit4)

# How do the models with lots of predictors perform?

for(x in 1:4) {
  p <- predict.lm(models[[x]], newdata = flips_test_data)
  MSE <- round(mean((flips_test_data$heads - p)^2), 3)
  print(paste("Test data MSE for model", x, "is", MSE))
}

# [1] "Test data MSE for model 1 is 0.25"
# [1] "Test data MSE for model 2 is 0.3"
# [1] "Test data MSE for model 3 is 0.544"
# [1] "Test data MSE for model 4 is 0.7"
```

Types of cross-validation

- Validation set approach
- Leave-one-out cross-validation (LOOCV)
- k -fold cross-validation

Validation sets

1. Randomly divide the data into two "folds"
 - One fold is used to **train** the model
 - Second fold is held out and later used to **test** the model
2. Fit the model on the training data
3. Use fitted model from step 2 to predict outcomes in the test data
4. Error rate is the MSE of the predictions in step 3

Key drawback: training data only includes half the observations

Leave-one-out cross-validation (LOOCV)

1. Choose 1 observation from the data to be the test data
 - Remaining $n - 1$ observations used to train the model
2. Fit the model on the training data
3. Use fitted model from step 2 to predict outcomes in test data (which has 1 data point)
4. Repeat steps 1-3 for each observation in the data
 - Error rate is the average MSE across all test datasets

Key drawback: computationally intensive, especially for large datasets (why?)

k -fold cross-validation offers a balance

1. Randomly divide the data into k folds
 - Select one fold to **test** the model
 - Remaining $k - 1$ folds are used to **train** the model
2. Fit the model on the training data ($k - 1$ folds)
3. Use fitted model to predict outcomes in the test data (1 fold)
4. Repeat steps 1-3, using a different fold to test the model each time
 - Error rate is the average MSE across all k folds

Note: LOOCV is special case of k -fold cross-validation with $k = n$

Which procedure works best for estimating test error?

- How to choose k , the number of folds?
- We face a bias-variance tradeoff in the cross-validated error estimate:
 - LOOCV ($k = n$)
 - Has more data in each training dataset → lower bias
 - But, those training datasets are nearly identical → higher variance
 - Small k (e.g., 2-fold)
 - Training sets are much smaller → CV error is biased upward
 - Error estimates are averaged over larger test folds → lower variance
- In practice, $k = 5$ or $k = 10$ is usually a good compromise, delivering the best estimate

Cross-validation in R

Try it: load libraries, turn off scientific notation

```
library(tidyverse)
```

```
library(ggplot2)
```

scipen: a penalty for deciding whether to print fixed or exponential notation.

Positive values encourage fixed notation, negative encourage scientific notation

```
options(scipen=999)
```

Try it: load and inspect the Boston housing dataset

```
# Housing information for 506 areas around Boston  
housing <- read_csv("lecture-10-housing.csv")  
nrow(housing)  
ncol(housing)  
  
# [1] 506  
# [1] 14
```

Variables come from the 1970 Census

Variable name	Definition
MEDV	Median value of owner-occupied homes in \$1000's
CRIM	Per capita crime rate by census tract
ZN	Proportion of residential land zoned for lots over 25,000 sq.ft.
INDUS	Proportion of non-retail business acres per town
CHAS	Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
NOX	Nitric oxides concentration (parts per 10 million)
RM	Average number of rooms per dwelling
AGE	Proportion of owner-occupied units built prior to 1940
DIS	Weighted distances to five Boston employment centers
RAD	Index of accessibility to radial highways
TAX	Full-value property-tax rate per \$10,000
PTRATIO	Pupil-teacher ratio by town
LSTAT	% lower status of the population

Model housing value as function of the number of rooms

- MEDV: median housing value in the area in 1970
- RM: average number of rooms per house in the area

```
summary(housing$MEDV)
cat("\n")
summary(housing$RM)
```

```
#      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
#      5.00   17.02  21.20    22.53  25.00    50.00
#
#      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
#      3.561   5.886  6.208    6.285  6.623    8.780
```

Employ a polynomial in the number of rooms

- We will model value as function of j -degree polynomial in the number of rooms

$$MEDV = \beta_0 + \sum_{i=1}^j \beta_i RM^i + \epsilon$$

- What degree (j) should we use?
- Use cross-validation to decide

Try it: validation set approach

```
# Set seed (why is this needed?)  
set.seed(1)  
n_rows <- nrow(housing)  
  
# Randomly select half the observations  
train <- sample(n_rows, n_rows/2)  
  
# RM degree 1  
# 1. Estimate model using training observations only by specifying 'subset' option  
lm_1 <- lm(MEDV ~ RM, data = housing, subset = train)  
  
# 2. Calculate MSE using test data  
lm_1_mse <- mean((housing$MEDV - predict(lm_1, housing))[-train]^2)  
lm_1_mse  
  
# [1] 49.65573
```

Validation set approach: higher degrees

```
# RM degree 2
lm_2 <- lm(MEDV ~ poly(RM, 2), data = housing, subset = train)
lm_2_mse <- mean((housing$MEDV - predict(lm_2, housing))[-train]^2)
lm_2_mse

# RM degree 3
lm_3 <- lm(MEDV ~ poly(RM, 3), data = housing, subset = train)
lm_3_mse <- mean((housing$MEDV - predict(lm_3, housing))[-train]^2)
lm_3_mse

# [1] 50.59916
# [1] 48.63831
```

This code is getting repetitive....

- Write a function!
- Function will calculate the MSE for different models
 - Different training samples
 - Different polynomial orders

Try it: write the function

```
f_mse <- function(order=10, seed=NULL) {  
  mse_lm <- rep(0,order)  
  set.seed(seed)  
  n_rows <- nrow(housing)  
  train <- sample(n_rows, n_rows/2)  
  
  for (j in 1:order){  
    # Estimate training model where Y = MEDV and X = polynomial in rooms of order j  
    lm <-  
  
    # Calculate the MSE using test data (hint: look at code on previous slide)  
    mse_lm[j] <-  
  }  
  
  return(mse_lm)  
}
```

Write the function

```
f_mse <- function(order=10, seed=NULL) {  
  mse_lm <- rep(0,order)  
  set.seed(seed)  
  n_rows <- nrow(housing)  
  train <- sample(n_rows, n_rows/2)  
  
  for (j in 1:order){  
    # Estimate training model where Y = MEDV and X = polynomial in rooms of order j  
    lm <- lm(MEDV ~ poly(RM, j), data = housing, subset = train)  
  
    # Calculate the MSE using test data (hint: look at code on previous slide)  
    mse_lm[j] <- mean((housing$MEDV - predict(lm, housing))[-train]^2)  
  }  
  
  return(mse_lm)  
}
```

Try it: call the function using different seeds

```
# Confirm that we get the same results as previous slide (RM degree 1, 2, 3)
f_mse(3,1)

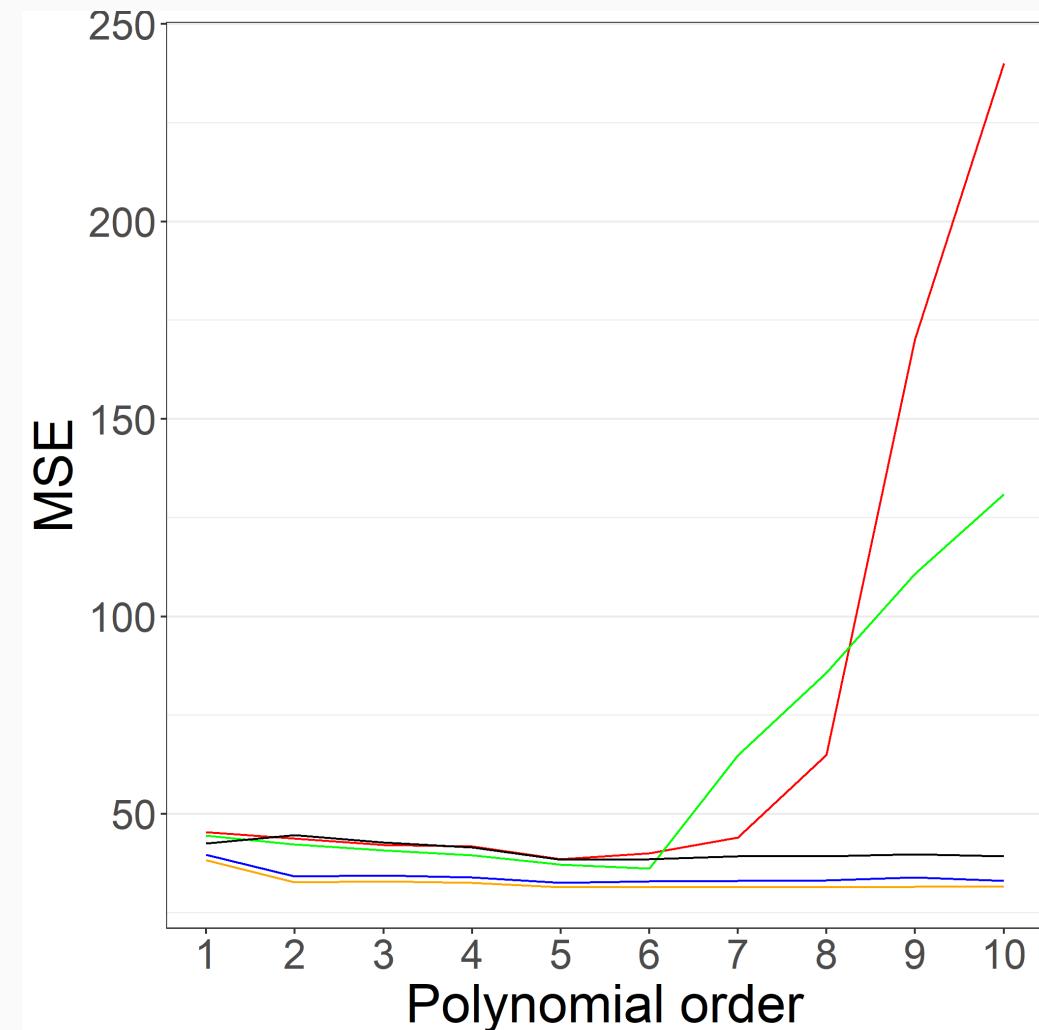
mse_1 <- f_mse(10, 891)
mse_2 <- f_mse(10, 892)
mse_3 <- f_mse(10, 893)
mse_4 <- f_mse(10, 894)
mse_5 <- f_mse(10, 895)

# Build a data frame with the results
mse <- data.frame(mse_1, mse_2, mse_3, mse_4, mse_5) %>%
  mutate(order = row_number())

# [1] 49.65573 50.59916 48.63831
```

Try it: plot our results

```
# Show MSE as a function of the polynomial degree
# for the 5 different seeds we tried
ggplot(mse, aes(x=order)) + theme_bw() +
  geom_line(aes(y = mse_1), color = "red") +
  geom_line(aes(y = mse_2), color = "blue") +
  geom_line(aes(y = mse_3), color = "green") +
  geom_line(aes(y = mse_4), color = "orange") +
  geom_line(aes(y = mse_5), color = "black") +
  scale_x_continuous(breaks = c(1:10)) +
  ylab("MSE") + xlab("Polynomial order") +
  theme(
    axis.title = element_text(size = 26),
    axis.text = element_text(size = 20),
    panel.grid.major.x = element_blank(),
    panel.grid.minor.x = element_blank()
  )
```



Leave-One-Out Cross-Validation (LOOCV)

```
library(boot) # Calculate LOOCV using cv.glm(), part of 'boot' library

# RM degree 1. No need to set a seed with LOOCV (why not?)
glm_1 <- glm(MEDV ~ RM, data = housing) # GLM default is regression
cv_err_1 <- cv.glm(housing, glm_1)      # cv.glm default is LOOCV

# Value of K (LOOCV is k-fold with k = n)
cv_err_1$K

# Cross-validation estimate of prediction error
# We care about the first number: raw prediction error (not adjusted)
cv_err_1$delta[1]

# [1] 506
# [1] 44.21666
```

Estimate LOOCV for multiple polynomials (slow)

```
for (j in 1:12) {  
  glm_j <- glm(MEDV ~ poly(RM,j), data = housing)  
  cv_err_j <- cv.glm(housing, glm_j)$delta[1]  
  print(paste("MSE for order",j,"is",round(cv_err_j,2)))  
}  
  
# [1] "MSE for order 1 is 44.22"  
# [1] "MSE for order 2 is 39.13"  
# [1] "MSE for order 3 is 38.24"  
# [1] "MSE for order 4 is 38.5"  
# [1] "MSE for order 5 is 36.15"  
# [1] "MSE for order 6 is 37.05"  
# [1] "MSE for order 7 is 37.67"  
# [1] "MSE for order 8 is 37.99"  
# [1] "MSE for order 9 is 67.3"  
# [1] "MSE for order 10 is 36.63"  
# [1] "MSE for order 11 is 80.48"  
# [1] "MSE for order 12 is 1898.7"
```

k-fold cross-validation

```
set.seed(1)

## RM degree 1
glm_1 <- glm(MEDV ~ RM, data = housing)

# 10-fold cross-validation
cv_err_k10_1 <- cv.glm(housing, glm_1, K = 10)
cv_err_k10_1$K

# Cross-validation estimate of prediction error
cv_err_k10_1$delta[1]

# [1] 10
# [1] 44.49213
```

Estimate k-fold cross-validation for multiple polynomials

```
for (j in 1:12) {  
  glm_j <- glm(MEDV ~ poly(RM,j), data = housing)  
  cv_err_j <- cv.glm(housing, glm_j, K=10)$delta[1]  
  print(paste("MSE for order", j, "is", round(cv_err_j,2)))  
}  
  
# [1] "MSE for order 1 is 44.16"  
# [1] "MSE for order 2 is 39.44"  
# [1] "MSE for order 3 is 38.02"  
# [1] "MSE for order 4 is 38.37"  
# [1] "MSE for order 5 is 36.29"  
# [1] "MSE for order 6 is 37.25"  
# [1] "MSE for order 7 is 38.07"  
# [1] "MSE for order 8 is 37.63"  
# [1] "MSE for order 9 is 69.27"  
# [1] "MSE for order 10 is 51.33"  
# [1] "MSE for order 11 is 108.76"  
# [1] "MSE for order 12 is 2319.54"
```

Compare validation set, LOOCV, and 10-fold CV

```
set.seed(2)
order <- 12
mse_vs     <- rep(0,order)
mse_loocv <- rep(0,order)
mse_k10    <- rep(0,order)

for (j in 1:order) {

  train <- sample(n_rows, n_rows/2)
  lm_vs <- lm(MEDV ~ poly(RM,j), data = housing, subset = train)
  mse_vs[j] <- mean((housing$MEDV - predict(lm_vs, housing))[-train]^2)

  lm <- glm(MEDV ~ poly(RM,j), data = housing)
  mse_loocv[j] <- cv.glm(housing, lm)$delta[1]
  mse_k10[j]   <- cv.glm(housing, lm, K = 10)$delta[1]
}
```

Table of comparisons

```
mat <- matrix(c(mse_vs, mse_loocv, mse_k10),  
               nrow = order, ncol = 3)  
rownames(mat) <- paste("Degree", rep(1:order))  
colnames(mat) <- c("Valid. Set", "LOOCV", "k-fold")  
kable(mat, digits=1, align="c")
```

- Which polynomial achieves the lowest MSE?
- Do validation set, LOOCV, and k -fold agree?

	Valid. Set	LOOCV	k-fold
Degree 1	39.8	44.2	44.0
Degree 2	34.3	39.1	39.3
Degree 3	45.7	38.2	38.8
Degree 4	38.7	38.5	41.6
Degree 5	40.1	36.1	37.3
Degree 6	1984.8	37.0	36.5
Degree 7	39.0	37.7	38.1
Degree 8	32.5	38.0	38.1
Degree 9	3149.9	67.3	77.5
Degree 10	69.9	36.6	36.6
Degree 11	52.8	80.5	96.1
Degree 12	11003.4	1898.7	2269.8

Summary

- Model fit needs to be assessed using a test dataset
- This assessment is accomplished using cross-validation
 - Validation sets
 - Leave-one-out cross-validation
 - k -fold cross-validation
- k -fold cross-validation with 5 or 10 folds generally performs best
- Lab-10 due Sunday at 11:59pm
- Make sure to **stop your instance** when you are done working