

# Lecture 13

## Regression and classification trees

---

Julian Reif  
Fall 2025

# RStudio setup for this lecture

- Log into RStudio on your Amazon EC2 instance
  - Use AMI **FIN550-RStudio** with IAM role **BigDataEC2Role**

*# Enter this command via RStudio Terminal*

```
aws s3 cp --recursive s3://bigdata-fin550-reif/lecture-13 ~/fin550/lecture-13
```

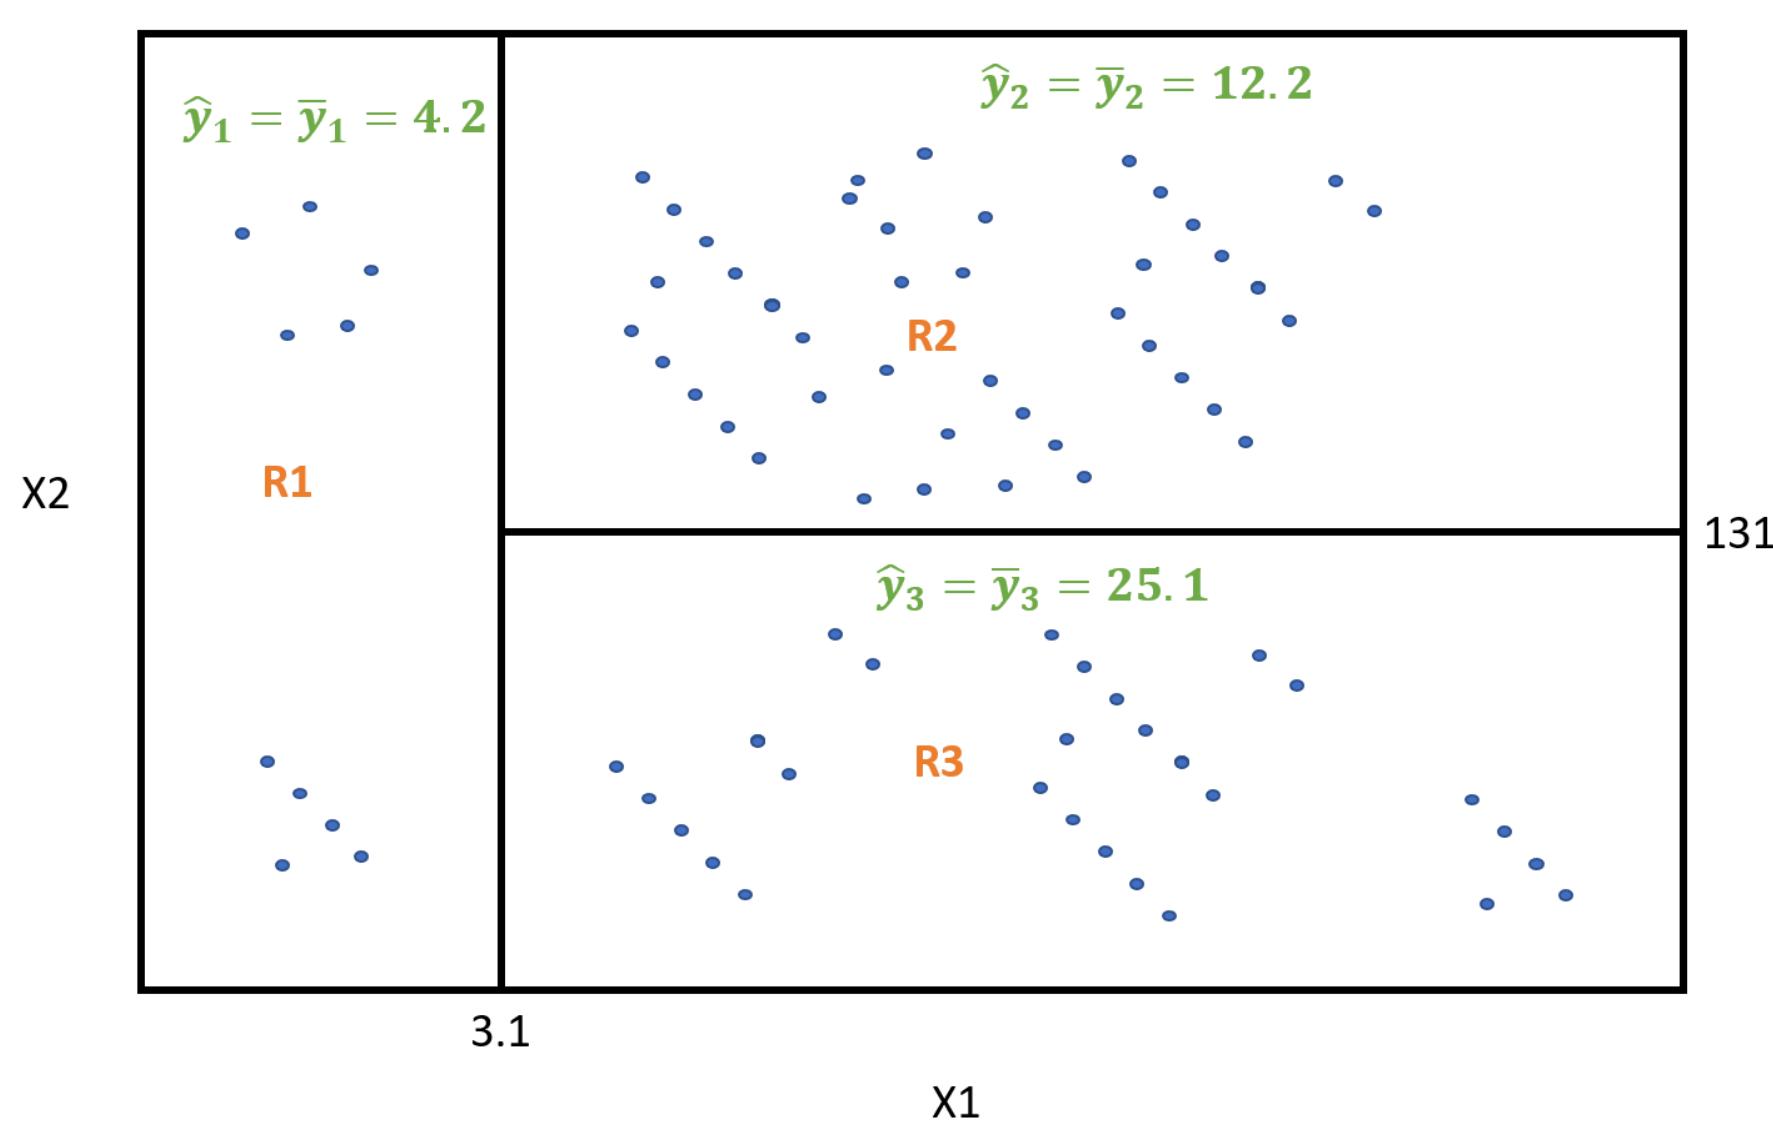
# Decision tree theory

---

# Decision trees

- Trees can be used for regression (continuous outcome) or classification analysis
- Building a tree consists of two main steps:
  1. Divide the predictor space up into different regions
  2. For every observation in a region, the prediction is the mean/mode of that region
    - Regression tree: mean
    - Classification tree: mode
- I will cover regression trees first

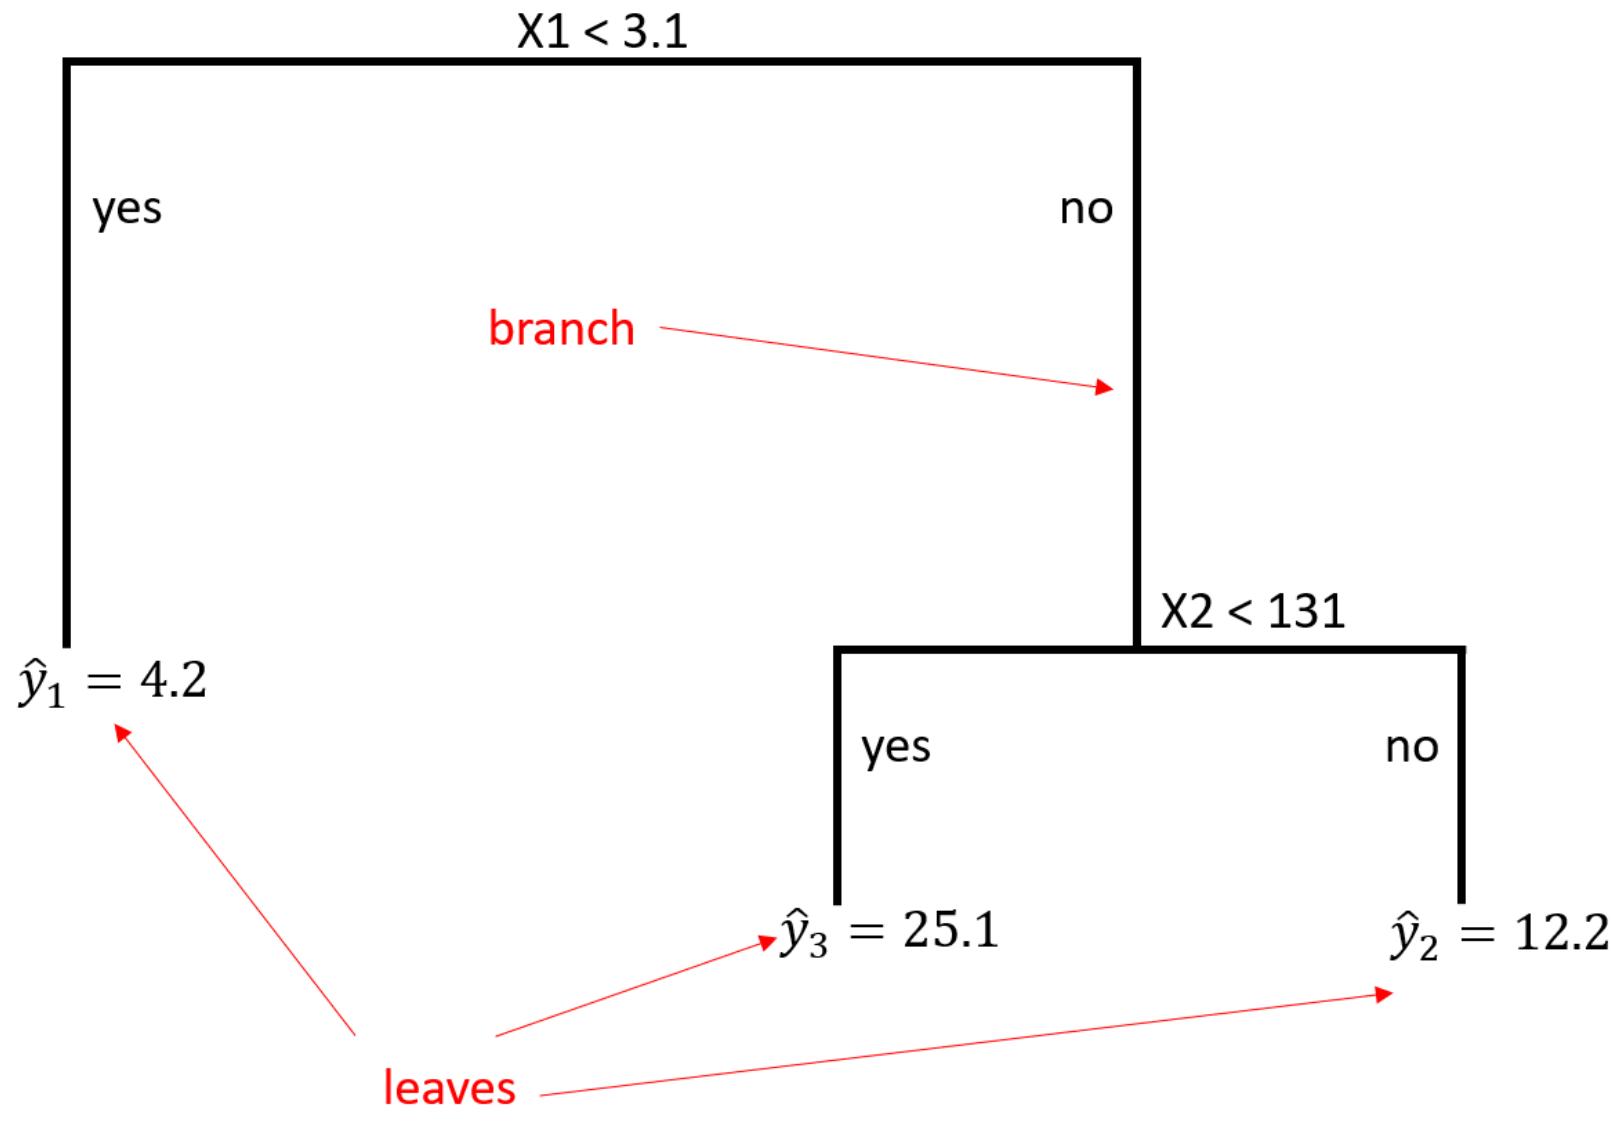
# Example: dividing predictor space into three regions



# The division is often represented in tree form

- Each region is called a "leaf" or a "terminal node"
- Connections between the internal nodes are called "branches"
- Size of tree,  $|T|$ , is equal to number of leaves ( $= \# \text{ splits} + 1$ )

# Example: regression tree with three leaves



# Use recursive binary splitting to grow a tree

- For each branch, we split the tree based on a predictor,  $X_j$ , and a value,  $s$
- This split will produce two distinct regions,  $R_1(j, s)$  and  $R_2(j, s)$ 
  - For example, if  $j = 2$  and  $s = 131$ , then we split based on  $X_2 < 131$ :
    - $R_1(2, 131)$  includes observations where  $X_2 < 131$
    - $R_2(2, 131)$  includes observations where  $X_2 \geq 131$
- We choose the split by solving for the  $j$  and  $s$  that minimize RSS:

$$RSS = \sum_{i:x_i \in R_1(j,s)} (y_i - \hat{y}_{R1})^2 + \sum_{i:x_i \in R_2(j,s)} (y_i - \hat{y}_{R2})^2$$

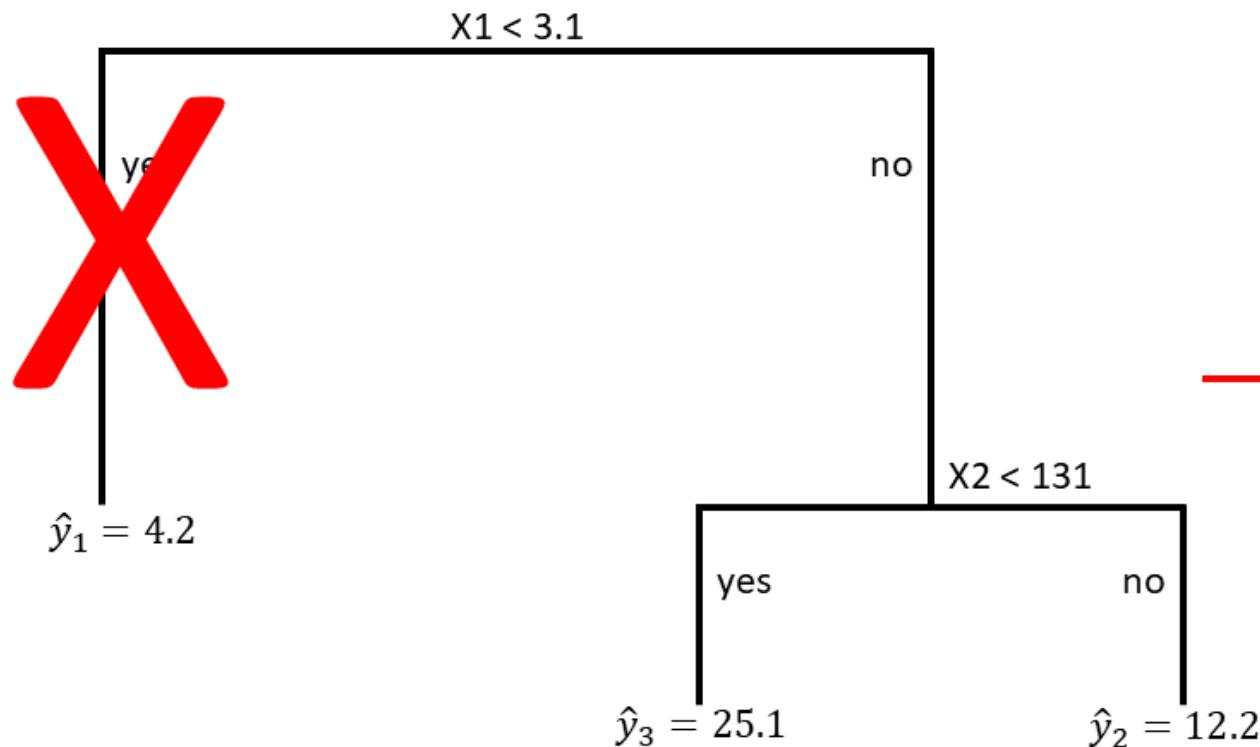
- Keep splitting until you reach a stopping criterion (e.g.,  $n < 5$  in each region)

# Pruning a tree

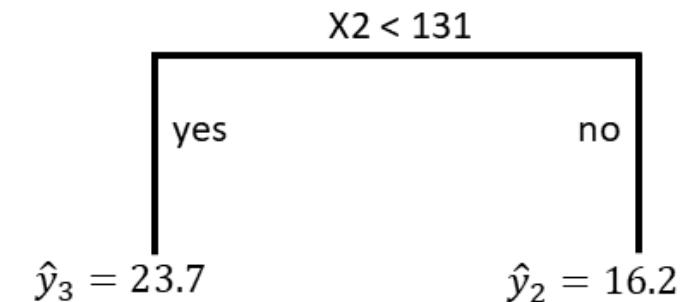
- Growing a tree by minimizing RSS usually leads to overfitting
  - Good fit to the training data
  - Poor fit to the test data
  - Note: linear regression has an analogous problem when there are too many variables
- Problem: some of the leaves are fitting to noise
- The solution is "tree pruning": dropping leaves with poor performance
  - Pruning produces a subtree

# Pruning a tree with T=3 leaves

Unpruned tree (T=3)



Pruned subtree (T=2)



# Pruning is accomplished by adding a penalty parameter

- Search through all possible subtrees, and find the one that minimizes:

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T| = RSS + \alpha |T|$$

where  $|T|$  is the number of leaves in subtree  $T$

- $\alpha \geq 0$  is a tuning parameter (aka "penalty" or "complexity" parameter)
  - When  $\alpha = 0$ , there is no pruning
  - As  $\alpha \rightarrow \infty$ , the number of leaves in the subtree goes to 1 (0 splits)
  - $\alpha$  is chosen using... cross-validation!

# Algorithm: building a regression tree

1. Grow a tree using recursive binary splitting
2. Consider a set of different  $\alpha$ 's:  $\alpha_1, \alpha_2, \dots, \alpha_k$
3. For each possible  $\alpha$ , prune the tree to find the optimal subtree
4. For each possible  $\alpha$ , calculate cross-validated mean-squared error for optimal subtree
  - Grow and prune a tree, using training data only
  - Calculate MSE using test data
5. Select the subtree from step 3 that corresponds to the "best"  $\alpha$  (smallest CV MSE)

# Classification trees are used with categorical outcomes

- In a regression tree, the prediction for a leaf is the **mean** of the observations
- In a classification tree, the prediction is the **mode**
- Example where  $y$  has  $K = 2$  categories and leaf has 3 obs: {yes, yes, no}
  - $\hat{y} = \text{"yes"}$  for these observations
- Example where  $y$  has  $K = 3$  categories and leaf has 5 obs: {MSF, MSF, MSF, MSFE, MSBA}
  - $\hat{y} = \text{"MSF"}$  for these observations
- Note: requires a tie-breaking rule for cases like {yes, yes, no, no}

# Classification trees uses node impurity instead of RSS

- We choose the split by solving for the  $j$  and  $s$  that minimizes "weighted Gini impurity":

$$\text{Impurity} = \frac{N_1(j, s)}{N} G_1(j, s) + \frac{N_2(j, s)}{N} G_2(j, s)$$

where  $N = N_1(j, s) + N_2(j, s)$  is total number of observations in leaves 1 and 2

- The impurity of node  $i$  is given by:

$$G_i = \sum_{k=1}^K \hat{p}_{ik} (1 - \hat{p}_{ik})$$

where  $\hat{p}_{ik}$  is the fraction of observations in leaf  $i$  belonging to class  $k$

# What is node impurity?

- Node impurity is a measure of classification error:

$$G_i = \sum_{k=1}^K \hat{p}_{ik} (1 - \hat{p}_{ik})$$

- Consider  $K = 2$  classes (yes/no), and let observations in the leaf be {yes, yes}
  - What is value of  $G$  for that leaf?

# What are the pros/cons of trees?

- Regression (and lasso) assume  $Y$  is a linear function of  $X$

$$Y = \beta_0 + \sum_{j=1}^p X_j \beta_j$$

- Regression trees, by contrast, allow for complicated nonlinearities:

$$Y = \sum_{m=1}^M c_m \mathbf{1}_{(X \in R_m)}$$

where  $\mathbf{1}$  is indicator function,  $R_m$  is a region, and  $c_m$  is the mean of that region

- However, in practice trees often perform worse than other methods
  - But, there are ways to improve predictive performance further (e.g., random forests)

# Regression trees in R

---

# Try it: load and inspect the Boston housing dataset

```
library(tidyverse)
library(ggplot2)

# Housing information for 506 areas around Boston
housing <- read_csv("lecture-13-housing.csv") %>%
  select(MEDV, CRIM, NOX, RM, TAX)
nrow(housing)
ncol(housing)

sum(!complete.cases(housing)) # Check for missing values
```

```
# [1] 506
# [1] 5
# [1] 0
```

# Variables come from the 1970 Census

Variable name	Definition
MEDV	Median value of owner-occupied homes in \$1000's
CRIM	Per capita crime rate by census tract
NOX	Nitric oxides concentration (parts per 10 million)
RM	Average number of rooms per dwelling
TAX	Full-value property-tax rate per \$10,000

```
summary(housing$MEDV)
```

```
#      Min. 1st Qu. Median     Mean 3rd Qu.    Max.
#      5.00   17.02  21.20   22.53  25.00   50.00
```

# Build decision trees using the rpart package

- Use `rpart` and `rpart.plot` libraries
- Key functions:
  - `mytree <- rpart()`: build the tree
  - `mytree$where`: where (which leaf) each observation was assigned to
  - `mytree$variable.importance`: list of most important predictors
  - `cp <- printcp(mytree)`: object containing tuning parameter values and CV MSE's
  - `prune(mytree, cp = myalpha)`: prune a tree for a particular value of alpha
  - `predict(mytree, newdata = mydata)`: use tree to form predictions
  - `prp(mytree)`: plot the tree

# Try it: build a regression tree

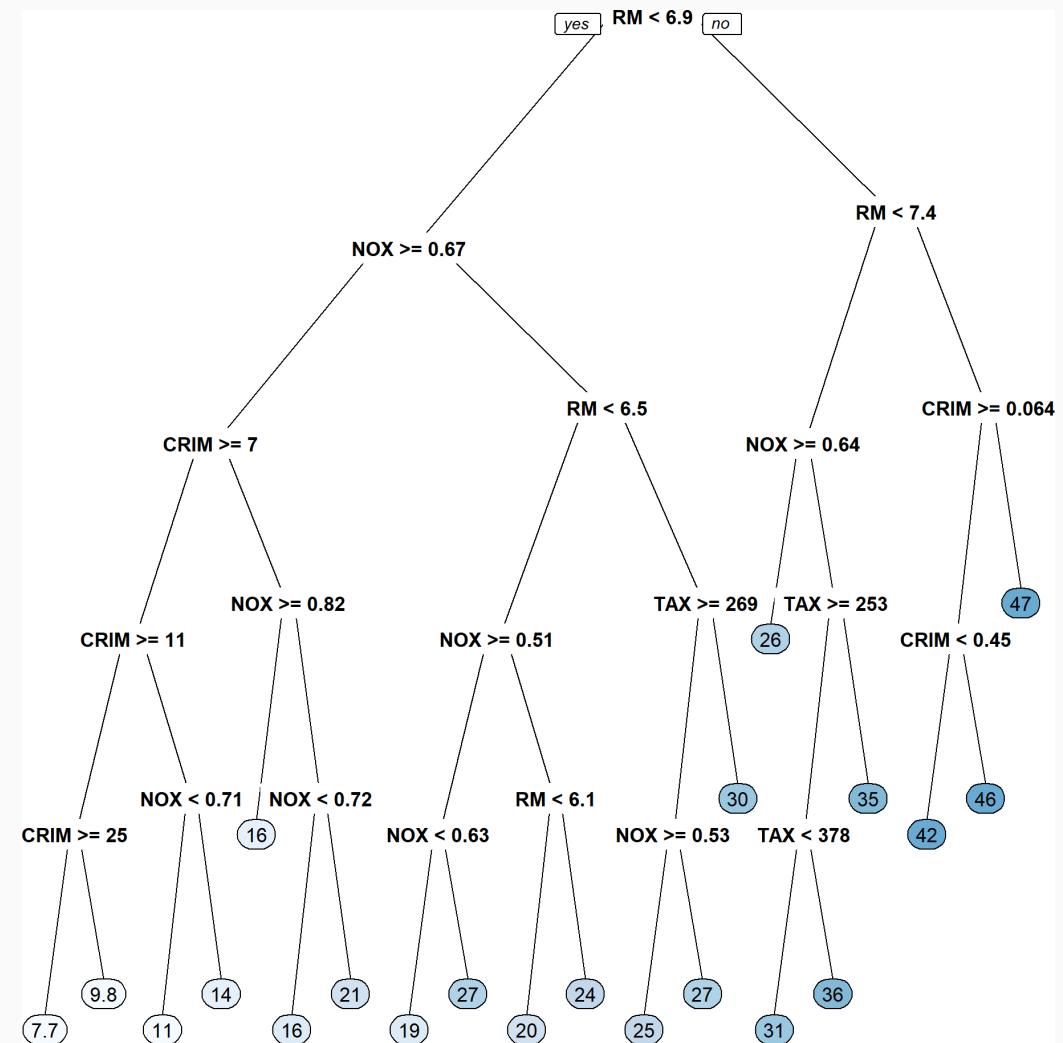
```
library(rpart)
library(rpart.plot)

# method="anova" for regression trees
# Set cp (alpha) equal to 0 to grow a full tree
# (Optional) set maxdepth=5 to limit how deep the tree is
set.seed(1)
tree_housing <- rpart(MEDV ~ ., data = housing, cp = 0,
                      maxdepth = 5, method = "anova")

# Plot the full tree
prp(tree_housing, box.palette = "auto")
```

# Plot the full tree (max depth = 5)

- What does tree look like if you omit  
maxdepth=5 from `rpart()`?



# Try it: number of leaves and variable importance

```
# Number of leaves/nodes  
n_distinct(tree_housing$where) # `where` reports which leaf each obs was assigned to  
  
# List variables that are most important for predicting outcomes  
# Do these variables lie near the top or the bottom of the tree?  
round(tree_housing$variable.importance,1)  
  
# [1] 21  
#      RM       NOX      CRIM      TAX  
# 25867.4  7704.5  5017.5  2968.9
```

# Try it: identify the vector of alpha's

```
cp <- printcp(tree_housing)

head(cp) # CP is alpha, xerror is CV MSE. What value is nsplit for the first CP? Why?
cat("\n")
as.vector(cp[,1]) # set of alpha's considered by `rpart`
```

```
#          CP nsplit rel error      xerror       xstd
# 1 0.45274420      0 1.0000000 1.0050753 0.08301497
# 2 0.11251782      1 0.5472558 0.5741582 0.05404496
# 3 0.07165784      2 0.4347380 0.4715515 0.06061093
# 4 0.05403671      3 0.3630801 0.4205866 0.05868427
# 5 0.02598598      4 0.3090434 0.3434092 0.05592466
# 6 0.01641021      5 0.2830575 0.3211350 0.05572784
#
# [1] 0.4527442007 0.1125178160 0.0716578409 0.0540367050 0.0259859832
# [6] 0.0164102099 0.0076138107 0.0073995566 0.0044752754 0.0042760000
# [11] 0.0028166156 0.0020539519 0.0019526403 0.0016295590 0.0012018315
# [16] 0.0006250947 0.0000000000
```

# Try it: identify the min CV MSE and the optimal alpha

```
# Use cp object to identify minimum CV MSE
cvmse.min <-

# Use cp object to identify optimal alpha (hint: use `which.min()` to find the index)
index <-
alpha.best <-

print(paste("Min CV MSE is", round(cvmse.min,3)))
print(paste("Optimal alpha is", round(alpha.best,5)))
```

# Identify the min CV MSE and the optimal alpha

```
# Use cp object to identify minimum CV MSE
cvmse.min <- min(cp[ , "xerror"])

# Use cp object to identify optimal alpha (hint: use `which.min()` to find the index)
index <- which.min(cp[ , "xerror"])
alpha.best <- cp[index, "CP"]

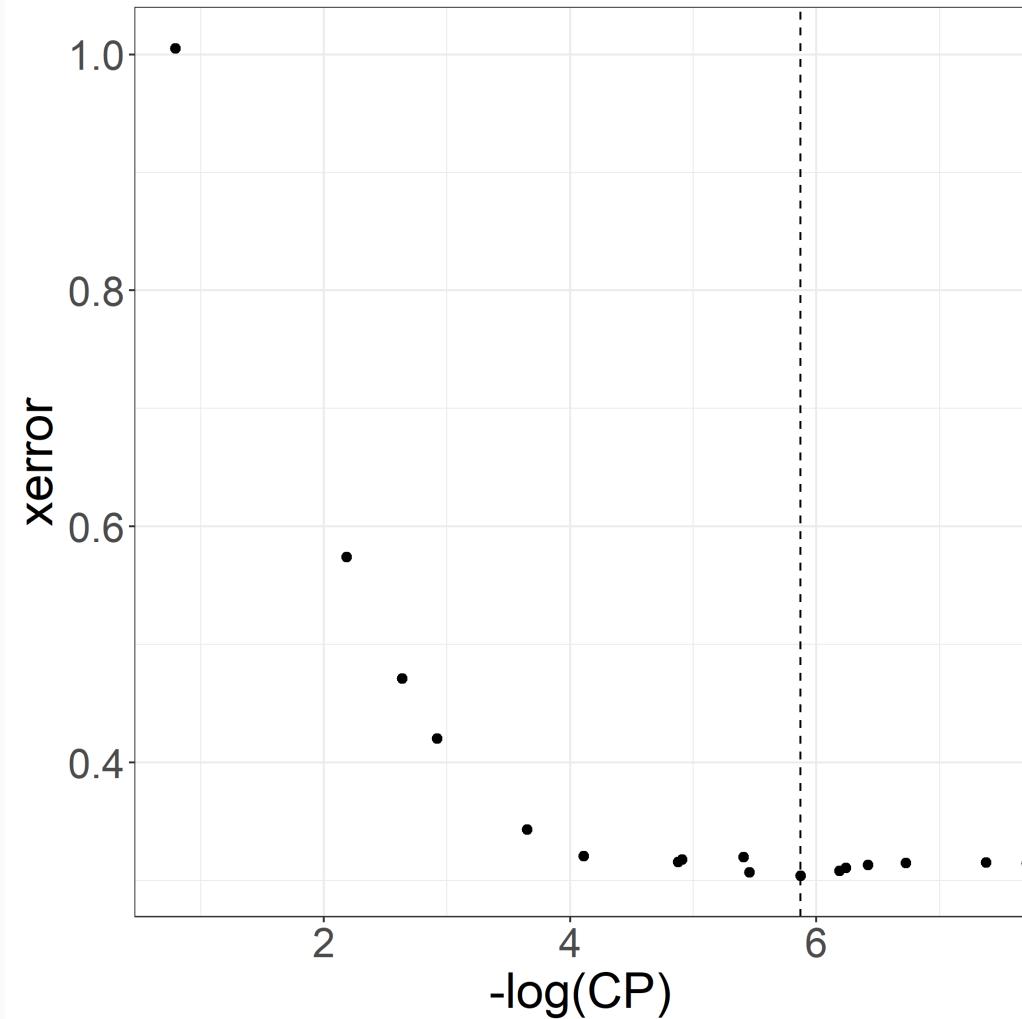
print(paste("Min CV MSE is", round(cvmse.min,3)))
print(paste("Optimal alpha is", round(alpha.best,5)))

# [1] "Min CV MSE is 0.304"
# [1] "Optimal alpha is 0.00282"
```

# Try it: plot CV MSE vs alpha

```
df <- as.data.frame(cp)
ggplot(df, aes(x=-log(CP), y=xerror)) +
  geom_point(size=2) +
  theme_bw() +
  theme(axis.title=element_text(size=24),
        axis.text=element_text(size=20))+
  geom_vline(xintercept =
              -log(alpha.best),
             linetype = "dashed")
-log(alpha.best)

# [1] 5.872219
```



# Try it: prune the tree

```
prune_housing <- prune(tree_housing, cp = alpha.best)
```

*# What happened to the number of leaves?*

```
n_distinct(tree_housing$where)
```

```
n_distinct(prune_housing$where)
```

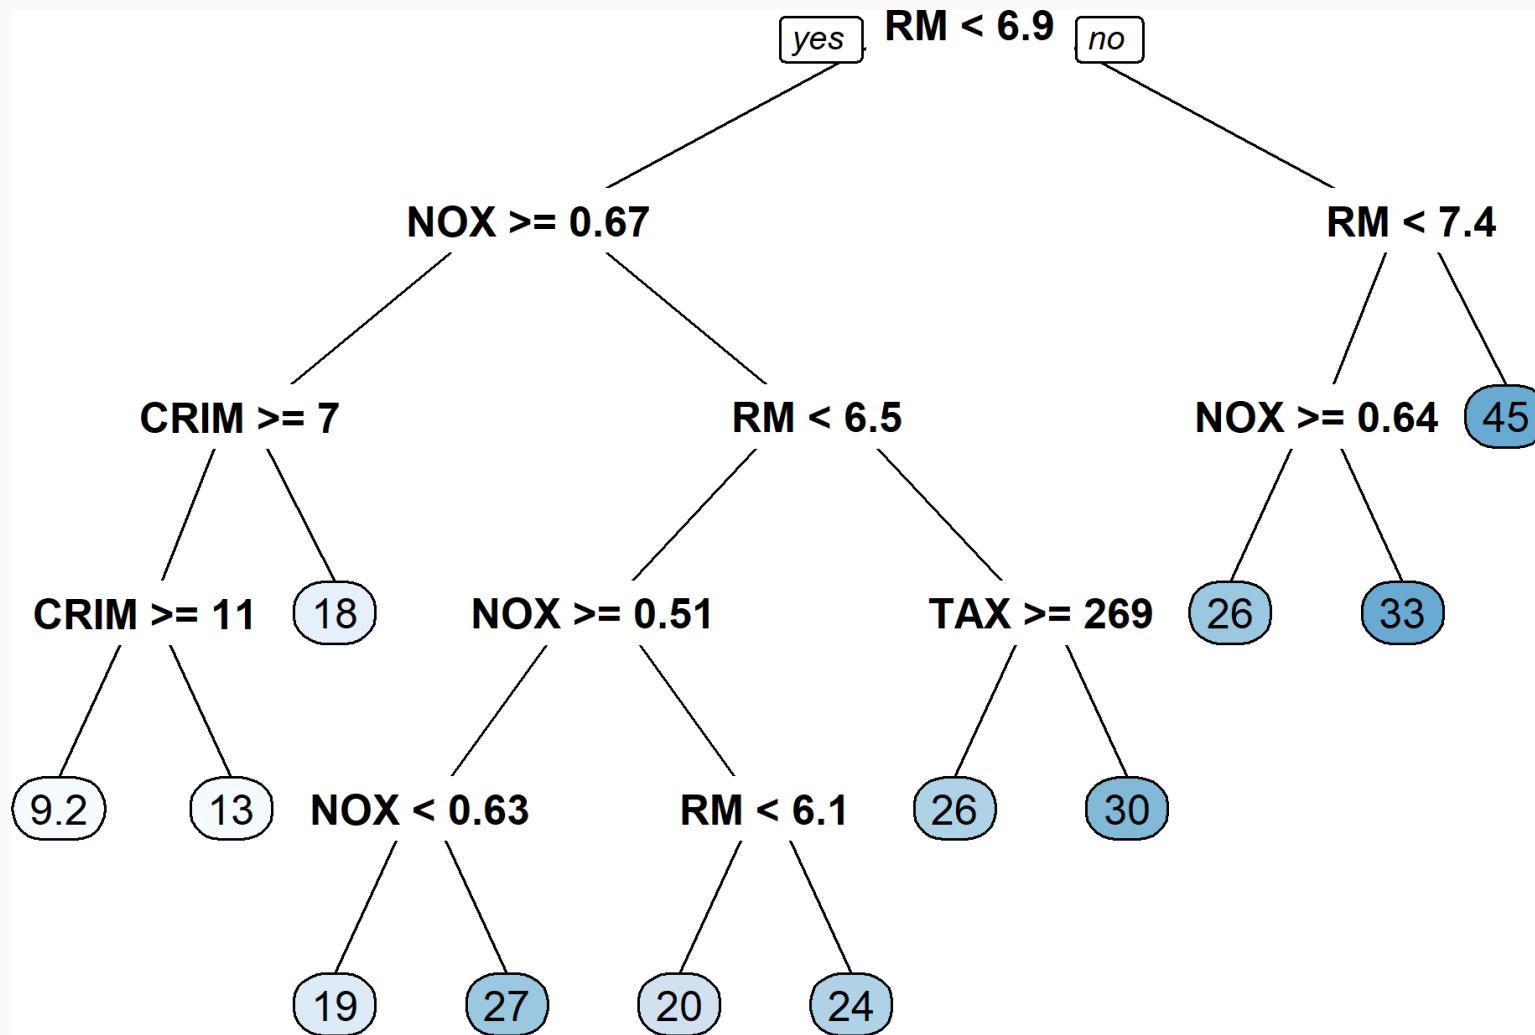
```
# [1] 21
```

```
# [1] 12
```

*# Plot the pruned tree*

```
prp(prune_housing, box.palette = "auto")
```

# Plot the pruned tree



# Make predictions

```
# Create a vector of predicted housing values  
medv_predict <- predict(prune_housing, newdata = housing)  
length(medv_predict)  
nrow(housing)  
  
# MSE  
mean((medv_predict - housing$MEDV)^2)  
  
# [1] 506  
# [1] 506  
# [1] 19.1187
```

# Summary

- Trees can be used for classification or regression problems
  - They are useful in presence of complicated nonlinearities
  - Otherwise, their predictive performance is often worse than other methods
- 
- Lab-13 due Sunday at 11:59pm
  - Reminder: midterm is **Wednesday, October 29**