

Lecture 14

Bagging, random forest, and boosting

Julian Reif
Fall 2025

RStudio setup for this lecture

- Log into RStudio on your Amazon EC2 instance
 - Use AMI **FIN550-RStudio** with IAM role **BigDataEC2Role**

Enter this command via RStudio Terminal

```
aws s3 cp --recursive s3://bigdata-fin550-reif/lecture-14 ~/fin550/lecture-14
```

Theory

What's wrong with decision trees?

- Decision trees are easy to explain and interpret
- But, they have high variance, delivering poor predictive performance
 - Pruning helps, but not enough
- Today, we discuss three improvements to standard decision trees:
 1. Bagging (special case of random forest)
 2. Random forest
 3. Boosting

Statistics review: bootstrapping

Bootstrapping lets you create more datasets

- Increasing data is useful for both training and testing
- How can you create new datasets without searching for more data?
- Bootstrap: draw randomly **with replacement** from your original dataset
 - Bootstrapped datasets all come from the same data generating process

Bootstrapping example

```
set.seed(2)
df <- data.frame(X1 = c(1,2,3,4,5,6,7))

df$X1_boot1 <- sample(df$X1, replace = T)
df$X1_boot2 <- sample(df$X1, replace = T)
df$X1_boot3 <- sample(df$X1, replace = T)
df
```

```
#   X1 X1_boot1 X1_boot2 X1_boot3
# 1  1       5       4       6
# 2  2       7       5       2
# 3  3       6       1       3
# 4  4       6       2       7
# 5  5       1       3       7
# 6  6       5       1       7
# 7  7       1       3       1
```

Bootstrapping properties

- Bootstrapped datasets are created by drawing **with replacement**
 - Some observations get drawn more than once
 - Some observations are not drawn at all
- What is the probability that an observation will not be in a bootstrapped dataset?

$$(1 - 1/n)^n \rightarrow 1/e \approx 0.368 \text{ when } n \text{ is large}$$

- Thus, about 1/3 of original observations will not be in a bootstrapped dataset

Bagging and random forests

Bootstrap aggregation ("bagging")

- Goal: reduce the variance of decision tree predictions
- If we have more datasets, we can reduce variance by averaging across them
 - Intuition: estimate of an average is better when $n = 1000$ than when $n = 100$
- Bagging uses bootstrap to create more datasets
 - Grow a separate tree on each dataset

Algorithm for bagging

1. Use bootstrap to create B datasets
2. For each dataset $b = 1, 2, \dots, B$, grow a full tree (no pruning)
 - Denote each tree's prediction as $\hat{f}^b(x)$
3. Calculate the average prediction across all the regression trees

$$\hat{f}_{bag}^b(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x)$$

- For classification tree: take a majority vote (i.e., mode of all $\hat{f}^b(x)$))
- Note: $B = 100$ is a good default, but it is fine to use larger values

Out-of-bag error estimation

- Recall: bootstrapping draws with replacement from the data
 - On average, only $2/3$ of the observations are used for tree b
- For each observation, form prediction **using trees grown without that observation**
 - On average, there are $B/3$ trees that didn't use this observation
- Take the average (or mode) of those $\approx B/3$ predictions
- "Out-of-bag" error is the mean-squared prediction error from these predictions


Random forests

- Random forests are identical to bagging, except for one difference
- When considering a split, use **random** sample of m predictors instead of all p predictors
- Example: there are 10 predictors, $X_1 \dots X_{10}$, and we set $m = 2$
 - Bagging: for all splits, search for optimal X_j and split point s among all 10 predictors
 - Random forest: consider X_1 and X_7 for first split, X_4 and X_5 for second split, etc.
- Typical default: $m = \sqrt{p}$
- Note: when $m = p$, random forest is same as bagging

Why limit the set of predictors??

- Without limiting the predictor set, the B trees used in bagging will be very similar
- Randomly dropping predictors limits the correlation among the B trees
 - Reducing the correlation reduces the variance of the prediction
- Random forest is "bagging + decorrelation"

Random forest generally outperforms bagging



Boosting

Boosting builds trees incrementally





- Build a series of **small**, incremental trees
 - Each tree is grown on a **dataset of residuals** from previous tree: $r_i = y_i - \hat{y}_i$
 - No bootstrapping
- Boosting depends on three tuning parameters:
 1. The number of trees, B
 - Chosen by cross-validation (note: very different from random forest!)
 2. The number of splits, d , in each incremental tree
 - Typically, we set $d = 1$ and consider a random sample of predictors
 3. The shrinkage parameter, λ
 - Controls the rate of learning from each incremental tree
 - Typically, $\lambda = 0.01$ or $\lambda = 0.001$

Algorithm for boosting


1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i
2. For $b = 1, 2, \dots, B$:
 - Fit a tree \hat{f}^b with d splits to the data (X, r)
 - Update \hat{f} using the learning parameter: $\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$
 - Update the residuals: $r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$
3. Calculate the boosted prediction across all the incremental trees:

$$\hat{f}_{boost}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$

Example: Boosting predictions for B=3, lambda = 0.5, d = 1



Boosting can outperform random forest



Bagging and random forest in R

Try it: load and inspect the Boston housing dataset

```
library(tidyverse)
library(ggplot2)

# Housing information for 506 areas around Boston
housing <- read_csv("lecture-14-housing.csv")
nrow(housing)
ncol(housing)

# The data are complete, so no need to remove observations
sum(!complete.cases(housing))
```

```
# [1] 506
# [1] 13
# [1] 0
```

Variables come from the 1970 Census

Variable name	Definition
MEDV	Median value of owner-occupied homes in \$1000's
CRIM	Per capita crime rate by census tract
ZN	Proportion of residential land zoned for lots over 25,000 sq.ft.
INDUS	Proportion of non-retail business acres per town
CHAS	Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
NOX	Nitric oxides concentration (parts per 10 million)
RM	Average number of rooms per dwelling
AGE	Proportion of owner-occupied units built prior to 1940
DIS	Weighted distances to five Boston employment centers
RAD	Index of accessibility to radial highways
TAX	Full-value property-tax rate per \$10,000
PTRATIO	Pupil-teacher ratio by town
LSTAT	% lower status of the population

Run random forest using the randomForest package

- Library `randomForest`
- Key functions:
 - `myforest <- randomForest()`: grow the forest
 - `myforest$ntree`: B , number of trees
 - `myforest$oob.times`: vector reporting # of times each observation is out-of-bag
 - `myforest$mse`: vector of mean-squared error for OOB observations for $b = 1\dots B$
 - `importance(myforest)` and `varImpPlot(myforest)`: most important predictors

Try it: build a model of housing value

```
summary(housing$MEDV)  
  
# How many predictors are available?  
num.p <-
```

Build a model of housing value

```
summary(housing$MEDV)

# How many predictors are available?
num.p <- length(names(housing))-1
cat("\n")
print(paste("Number of predictors is:", num.p))

#      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
# 5.00   17.02  21.20  22.53  25.00  50.00
#
# [1] "Number of predictors is: 12"
```

Try it: estimate a tree model using bagging

```
library(randomForest)
set.seed(1)

# Importance = T: calculate how important each variable is
# For bagging, number of splits m=p
bag_housing <- randomForest(MEDV ~ ., data = housing, mtry = num.p, importance = TRUE)
bag_housing

#
# Call:
# randomForest(formula = MEDV ~ ., data = housing, mtry = num.p,      importance = TRUE)
#           Type of random forest: regression
#                   Number of trees: 500
# No. of variables tried at each split: 12
#
#           Mean of squared residuals: 10.59575
#           % Var explained: 87.45
```

Try it: mean-squared error for "out-of-bag" observations

```
# B (number of trees)
bag_housing$ntree

# Number of times each observation is out-of-bag (OOB)
head(bag_housing$oob.times)

# On average, what fraction of the time is an observation OOB?
```

Mean-squared error for "out-of-bag" observations

```
# B (number of trees)
bag_housing$ntree

# Number of times each observation is out-of-bag (OOB)
head(bag_housing$oob.times)

# On average, what fraction of the time is an observation OOB?
cat("\n")
mean(bag_housing$oob.times)/500
1/(exp(1))
```

```
# [1] 500
# [1] 188 182 185 181 157 190
#
# [1] 0.3665968
# [1] 0.3678794
```


Average OOB error decreases with the number of trees

```
# Vector of OOB error (length=B)
round(head(bag_housing$mse), 1)

# [1] 53.8 32.1 27.7 26.4 22.9 19.1

# OOB error vs number of trees
plot(bag_housing)

# Same as:
#   mse <- bag_housing$mse
#   trees <- 1:bag_housing$ntree
#
#   df <- data.frame(x = trees, y=mse)
#   ggplot(df) + geom_line(aes(x=x, y=y))
```




Bagging: variable importance table

```
# Display in table form  
importance(bag_housing)
```

#	%IncMSE	IncNodePurity
# CRIM	17.320673	1781.32335
# ZN	4.265825	38.98644
# INDUS	9.246720	292.07306
# CHAS	3.555811	32.58706
# NOX	24.936067	1145.53829
# RM	59.780711	18530.57119
# AGE	19.086593	604.67534
# DIS	29.716657	2748.74198
# RAD	6.717248	158.17021
# TAX	13.555102	609.48993
# PTRATIO	18.574731	652.62634
# LSTAT	38.197511	15844.21364

Bagging: variable importance plot

```
# Or, visualize using a plot  
varImpPlot(bag_housing)
```



Try it: estimate a tree model using random forest

```
# Set number of random splits equal to a value less than 12
rf_housing <- randomForest(MEDV ~ ., data = housing, mtry = 6, importance = TRUE)
rf_housing


#
# Call:
# randomForest(formula = MEDV ~ ., data = housing, mtry = 6, importance = TRUE)
#           Type of random forest: regression
#           Number of trees: 500
# No. of variables tried at each split: 6
#
#           Mean of squared residuals: 9.59256
#           % Var explained: 88.64
```

Random forest: variable importance table

importance(rf_housing)	#	%IncMSE	IncNodePurity
	# CRIM	16.027037	2228.00483
	# ZN	3.094967	102.42133
	# INDUS	10.071582	1840.19628
	# CHAS	3.380017	88.20743
	# NOX	20.848299	2312.70179
	# RM	43.661166	15004.01671
	# AGE	12.403407	880.33134
	# DIS	22.138316	2656.43914
	# RAD	6.284539	193.00308
	# TAX	12.188322	798.85712
	# PTRATIO	15.650162	1895.00131
	# LSTAT	34.892093	14246.48393


Random forest: variable importance plot

```
# Or, visualize using a plot  
varImpPlot(rf_housing)
```



Our random forest model outperforms the bagging model

```
bag.mse <- bag_housing$mse  
rf.mse <- rf_housing$mse  
ntrees <- 1:bag_housing$ntree  
  
df <- data.frame(ntrees = ntrees,  
                  bag.mse = bag.mse,  
                  rf.mse = rf.mse)  
  
ggplot(df, aes(x=ntrees)) +  
  geom_line(aes(y=bag.mse),color="red") +  
  geom_line(aes(y=rf.mse),color="blue") +  
  ylab("Mean-squared Error") +  
  theme_bw() +  
  theme(axis.title=element_text(size=20))
```



Making predictions

```
# Note: for classification, specify type="prob" (probability) or type="class" (mode)
housing$yhat.bag <- predict(bag_housing, housing)
housing$yhat.rf <- predict(rf_housing, housing)

# MSE
mean((housing$yhat.bag - housing$MEDV)^2)
mean((housing$yhat.rf - housing$MEDV)^2)

# [1] 1.782202
# [1] 1.781661
```

Summary

- Random forest uses bootstrap to increase the number of trees you can grow
 - Averaging over these trees improves predictive performance
 - "Out-of-bag" observations provide a natural test dataset
- Bagging is a special case of random forest
- Boosting is another way to improve predictive performance of trees
- Lab-14 due Sunday at 11:59pm
- Reminder: midterm is **Wednesday, October 29**